

CS5250: Advanced Operating Systems

AY2021-2022 Semester 2

Weng-Fai Wong
COM2-03-56
wongwf@nus.edu.sg

Objectives of CS5250

- Understand the user needs and environmental constraints that drive the design of operating systems
- Understand the overall structure and design of any modern operating system especially in how their interact
- Understand the tradeoffs in the core algorithms and extend or modify them according to user needs
- Be equipped to take on more complex systems by being exposed to some of workings of actual systems

Syllabus

- Operating system design strategies including microkernels, mobile, embedded and real-time operating systems and the component's interfaces
- Priority and resource allocation strategies
- Scheduling algorithms
- Naming, protection and security
- User interface, windowing systems
- Distributed and shared objects
- File system implementations including network and distributed file systems
- Failure and recovery
- Virtualization and the Internet-ready OS

But how exactly will we do it?

- **The challenge: we can't write a complex OS from scratch**
 - Implementing the algorithms we will discuss often requires a lot of infrastructure support

- Deep dive into x86-Linux kernel source code
 - We will use Linux kernel 5.15.12

- Use hands-on assignment to advance our knowledge



The Linux Kernel Archives

About Contact us FAQ Releases Signatures Site news

Protocol Location
HTTP <https://www.kernel.org/pub/>
GIT <https://git.kernel.org/>
RSYNC <rsync://rsync.kernel.org/pub/>

Latest Release
5.15.12 [download icon]

mainline:	5.16-rc8	2022-01-02	[tarball]	[patch]	[inc. patch]	[view diff]	[browse]
stable:	5.15.12	2021-12-29	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	5.10.89	2021-12-29	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	5.4.169	2021-12-29	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	4.19.223	2021-12-29	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	4.14.260	2021-12-29	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	4.9.295	2021-12-29	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	4.4.297	2021-12-29	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
linux-next:	next-20211224	2021-12-24					[browse]


Prerequisites

- C/C++ programming
- X86 assembly programming (if possible)
- A prior undergraduate course in operating systems


Assessments

- 60% continual assessment
 - Pop quizzes
 - No tutorial
- 40% final assessment
 - Final assessment in the afternoon of **Friday, 6 May 2022**
- 4 take-home individual assignments


Reference Text




Supplementary texts (revision)



<https://www.amazon.sg/Modern-Operating-Systems-Andrew-Tanenbaum/dp/9332575770>



<https://www.amazon.sg/Operating-System-Concepts-Abraham-Silberschatz/dp/1118093755>



<https://www.amazon.sg/Operating-Systems-Internals-Design-Principles/dp/1292214295>

Also: <https://pages.cs.wisc.edu/~remzi/OSTEP/>

Introduction

What is an Operating Systems?

- Low level programs that sits between the hardware and user applications

Functions of an operating systems

- Manage a computer system's resources
- Provide user interface
- Provide services for applications

Key system resources

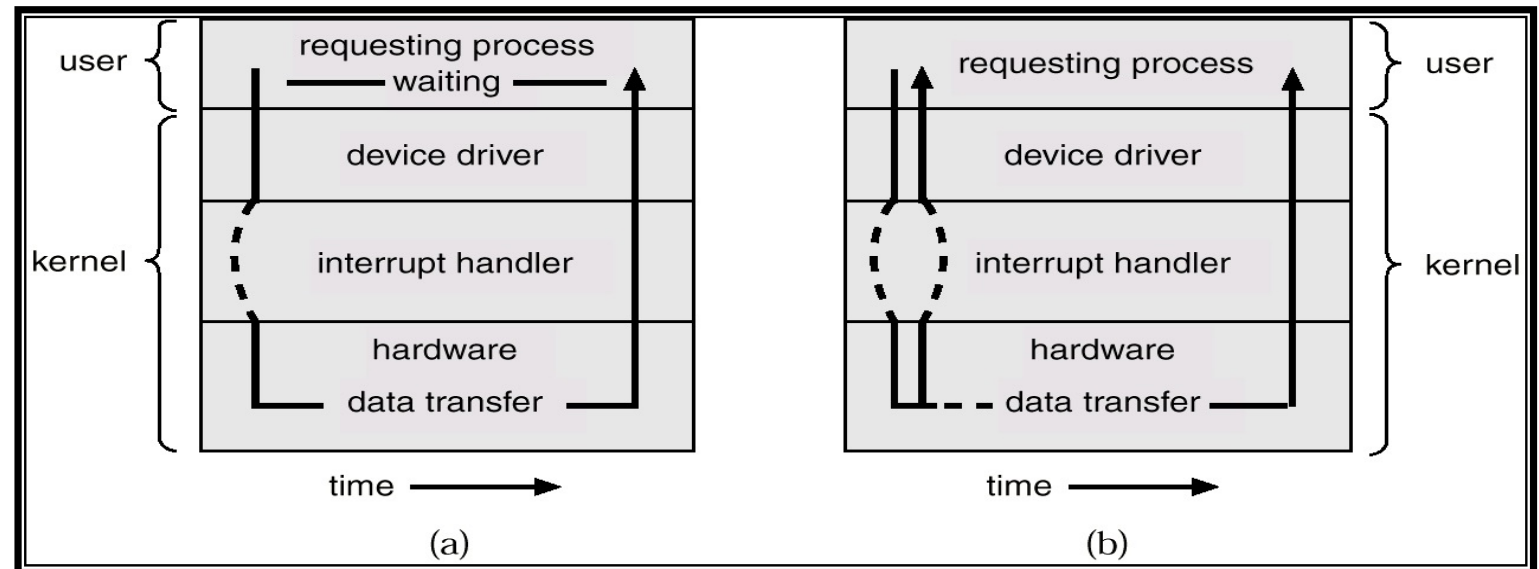
- Processors (cores)
- Memory
- Input/Output devices
 - Keyboard, mouse, monitor, and other human interface devices
 - Graphics card
 - Disk storage
 - Network
 - Printers, USB devices etc.
- Data and storage

Other key services

- Hardware abstraction
- Protection
- Sharing and data exchange
- Caching
- Interrupt handling

Hardware abstraction

- Provide regulated access to hardware
- Manufacturers provide “drivers” – software interfaces
- Provide book-keeping services



Protection


- Protect users (especially in a multiuser environment) against other users who may be malicious or plain ignorant
- At least two modes of protection enabled with hardware support:
 - User mode
 - Supervisor mode (root access)

Sharing

- Isolation vs sharing
- Need to establish that the sharing is legitimate
- Inter-process communication
- Via networking

Caching

- Caching is ubiquitous
 - To bridge differences in speeds and capacities
- Hardware caching
 - TLB, processor cache, disk caches
- Software caching
 - Page cache



Interrupt handling

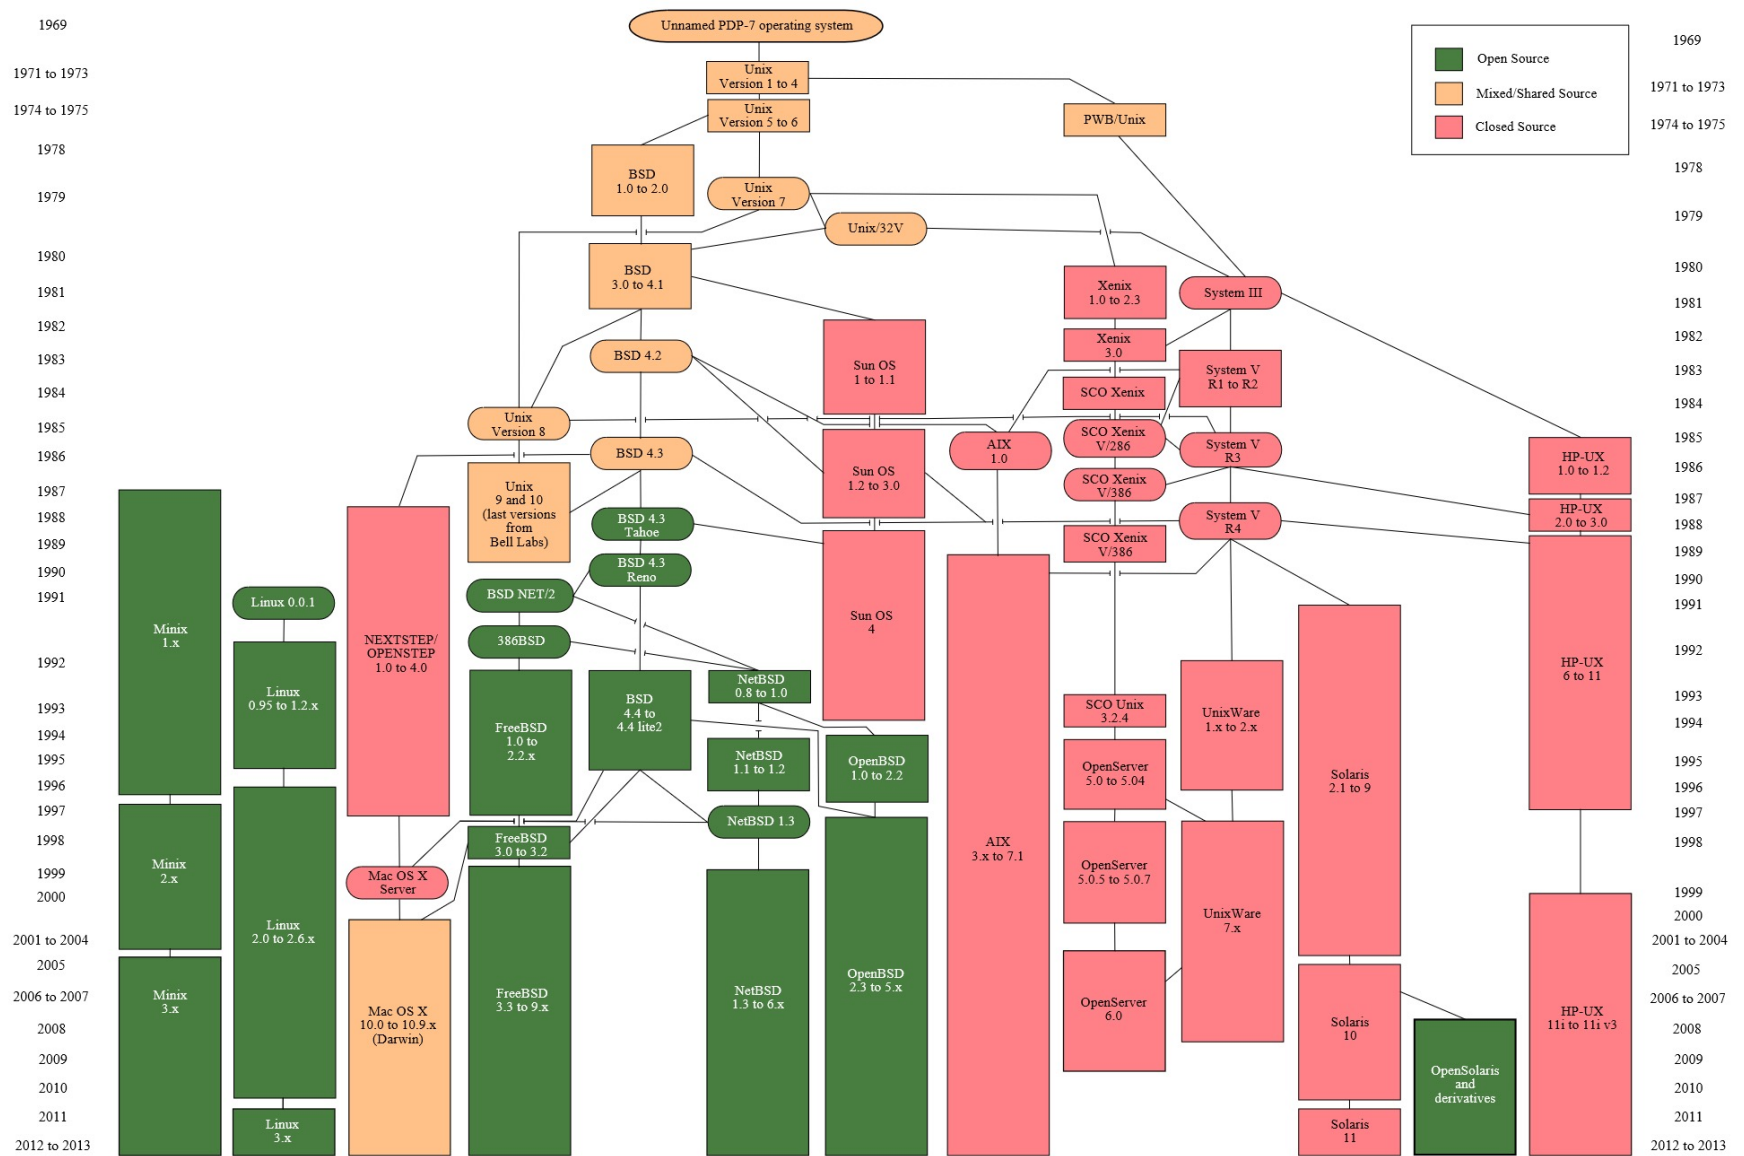
- Interrupts – asynchronous to program execution
 - caused by **external events**
 - may be handled **between** instructions, so can let the instructions currently active in the pipeline *complete* before passing control to the OS interrupt handler
 - simply suspend and resume user program
- Traps (Exception) – synchronous to program execution
 - caused by **internal events**
 - condition must be remedied by the trap handler for **that** instruction, so much stop the offending instruction *midstream* in the pipeline and pass control to the OS trap handler
 - the offending instruction may be retried (or simulated by the OS) and the program may continue or it may be aborted

Interrupt handling – examples

- Arithmetic overflow
- Undefined instruction
- TLB or page fault
- Segmentation fault
- I/O service request
- Hardware malfunction
- Timer

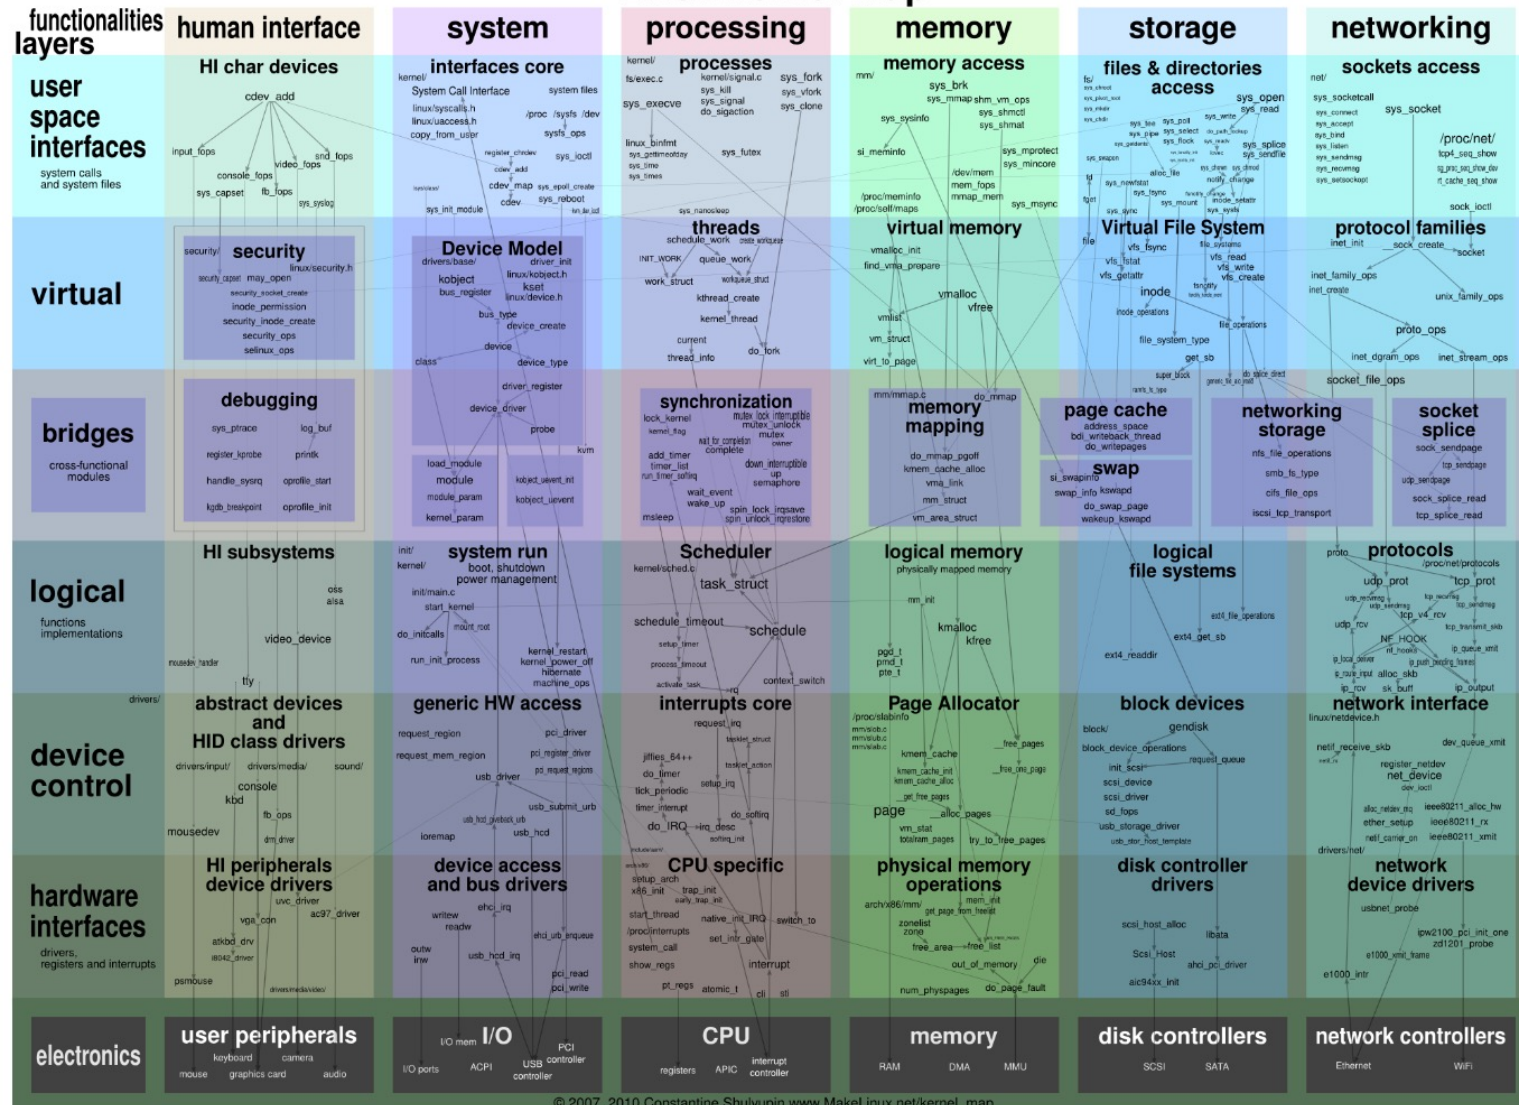
Key Operating Systems

Unix



Linux

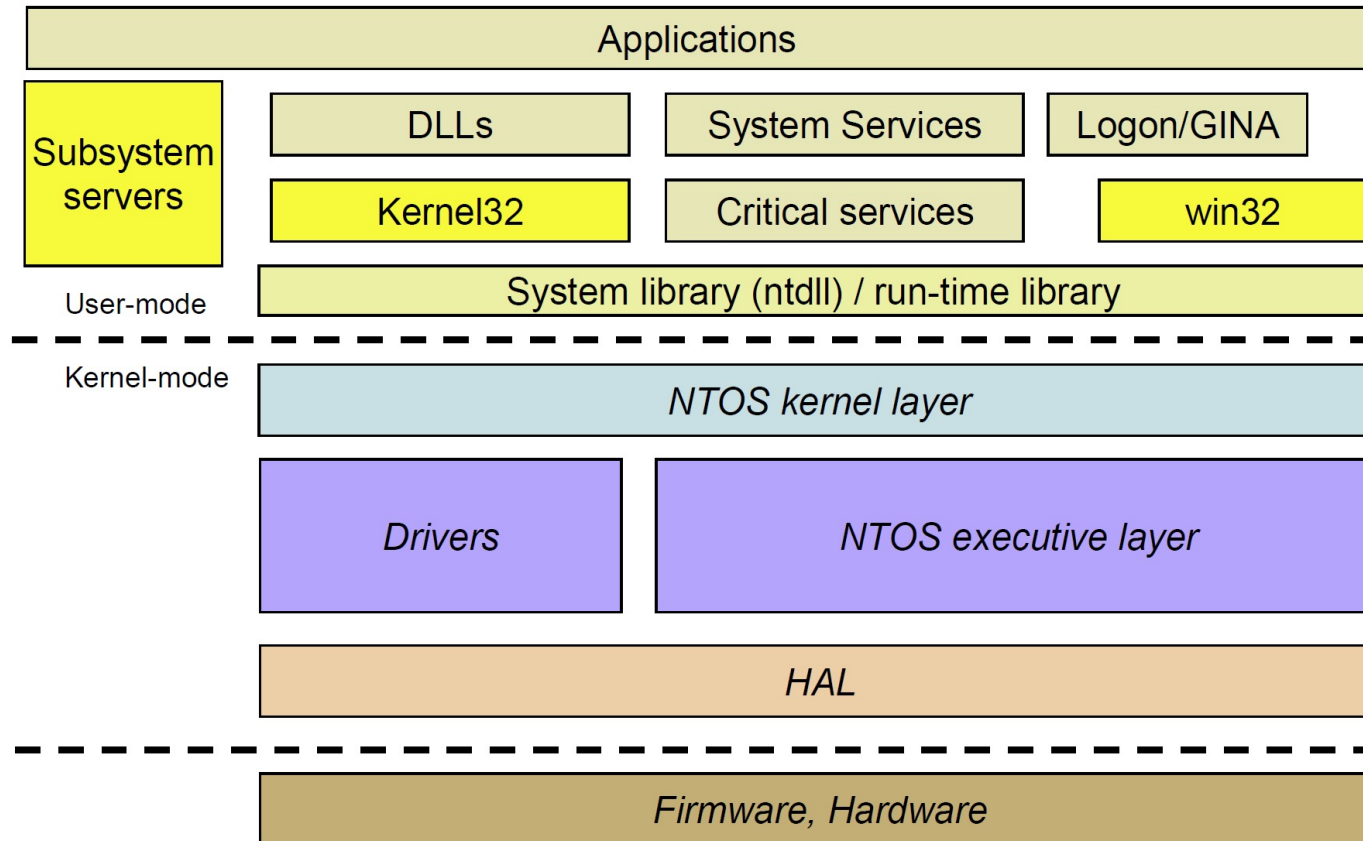
Linux kernel map



Microsoft Windows

- 16-bit:
 - MS-DOS (various versions, v1.0 in 1985)
 - Windows 1.X – 3.X, Windows 9X, Windows ME (2000)
- 32-bit:
 - Windows NT (32-bit, v3.1 in 1994)
 - Windows 2000, XP, 2003, Vista, 7, 8, 10 (2015)
- 64-bit:
 - Windows XP (2005), Vista, 7, 8, 10 (2015)
- Mostly on PC (Intel Processors) platforms
- Proprietary
 - some sources available under conditions
- Complex architecture, internals info not widely available


Microsoft Windows



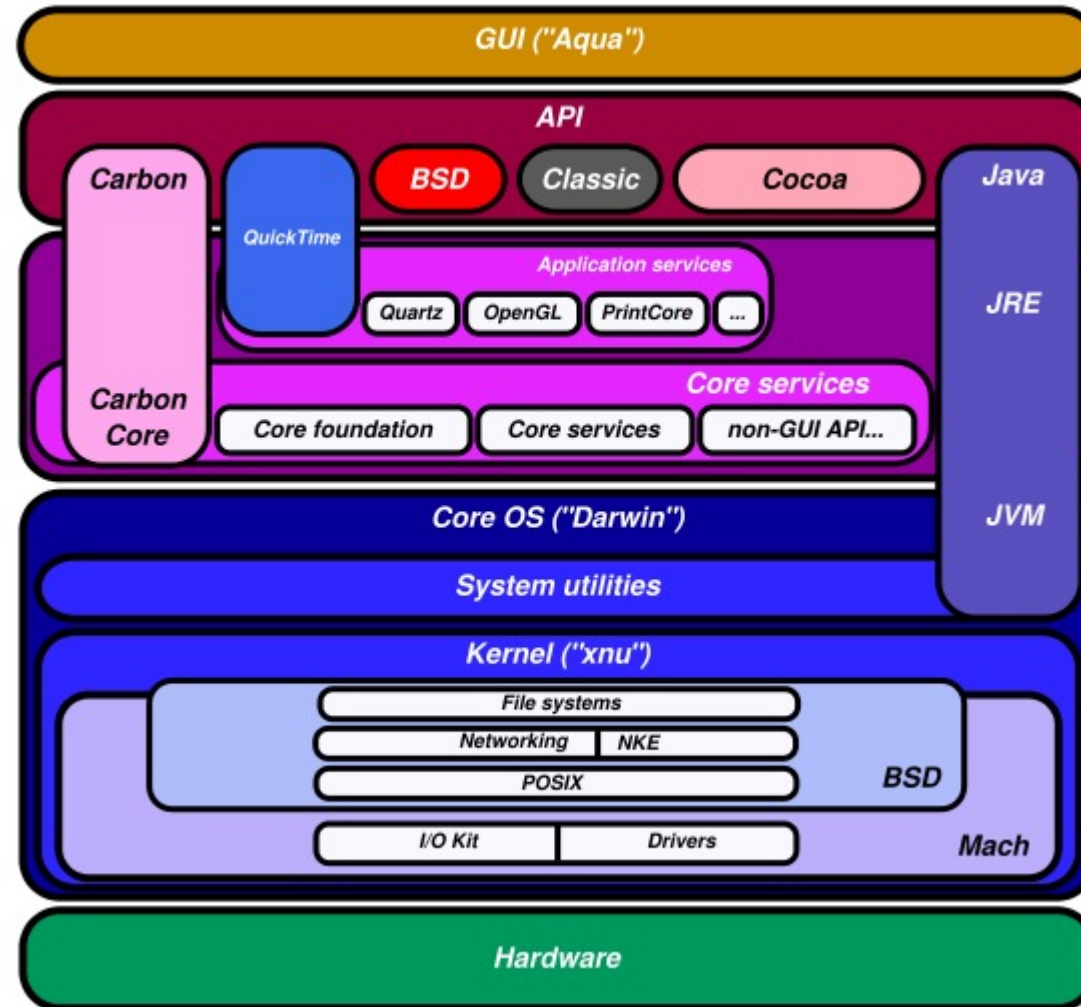
© Microsoft Corporation

MacOS

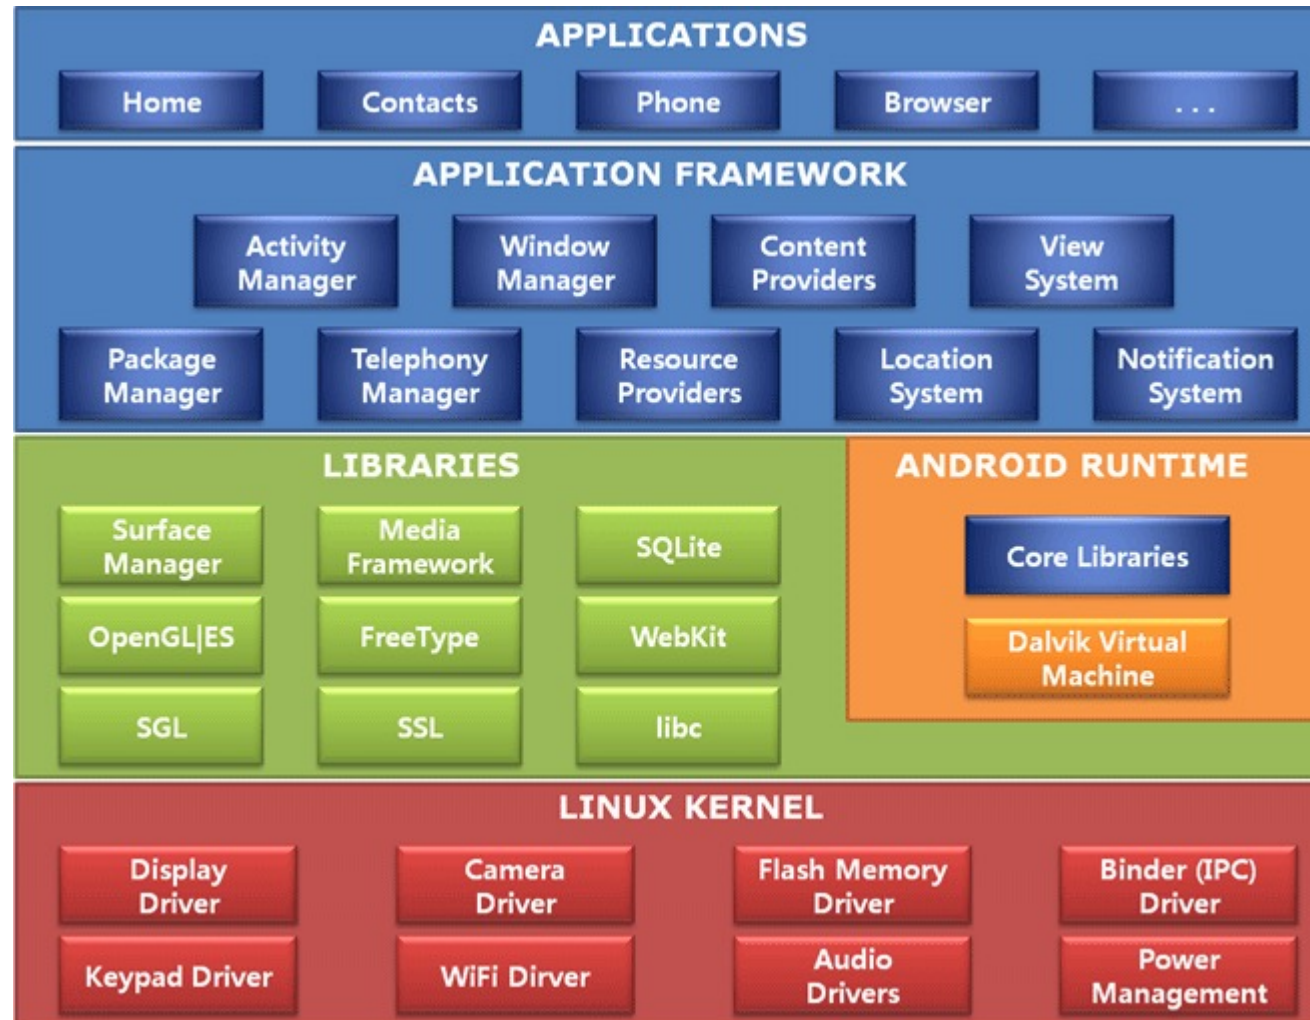
Timeline of Macintosh operating systems




MacOS



Android OS



The Show Stopper



```
A problem has been detected and windows has been shutdown to prevent damage to your computer.
If this is the first time you've seen this stop error screen, restart your computer. If this screen appears again, follow
these steps:
Check to make sure any new hardware or software is properly installed. If this is a new installation, ask your hardware
or software manufacturer for any windows updates you might need.
If problems continue, disable or remove any newly installed hardware or software. Disable BIOS memory options such as
caching or shadowing. If you need to use Safe Mode to remove or disable components, restart your computer, press F8 to
select Advanced Startup Options, and then select Safe Mode.
Technical information:
*** gv3.sys - Address F86B5A89 base at F86B5000, DateStamp 3dd9919eb
Beginning dump of physical memory
Physical memory dump complete.
Contact your system administrator or technical support group for further assistance.
```



Characteristics of OSes


Types of OS

- Batch
 - Data centers
- Interactive
 - PCs
 - Mobile devices
- Real-time
 - Embedded systems
 - Robots
- Hybrid

Structures of OS

- Monolithic
- Microkernel
- Network OS
- Distributed OS
- Exokernels

A Generic OS

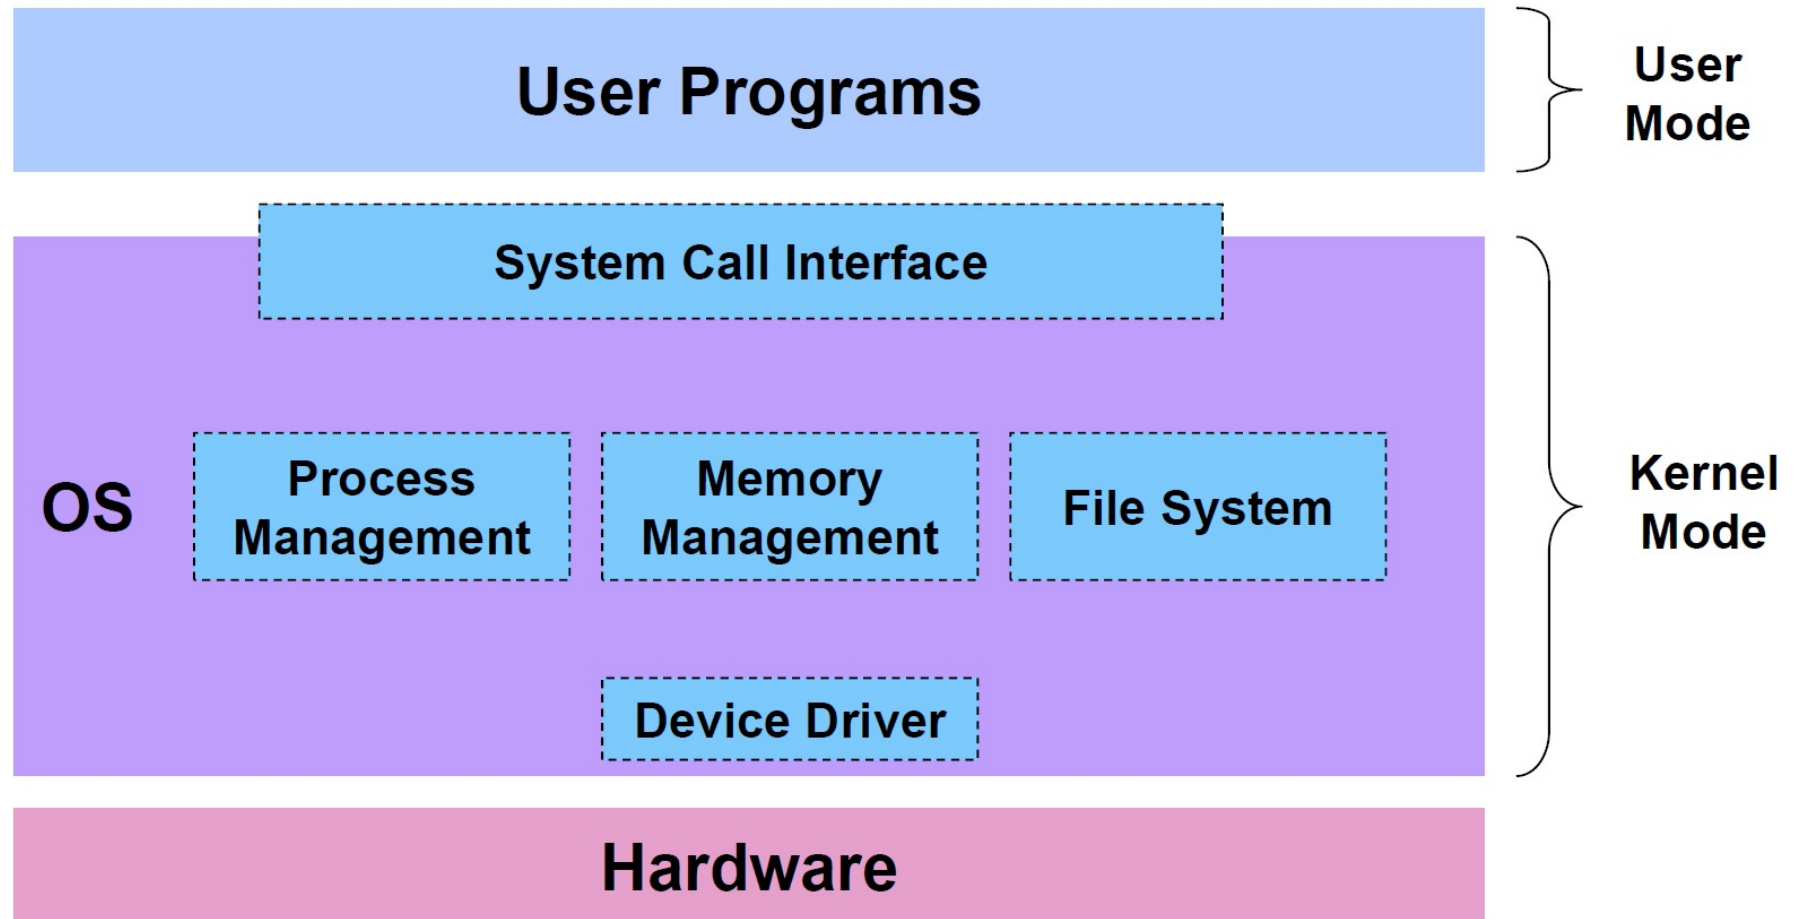


- **A**: OS executing machine instructions
- **B**: normal machine instructions executed (program/library code)
- **C**: calling OS using **system call interface**
- **D**: user program calls library code
- **E**: system processes
 - Provide high level services, usually part of OS

Monolithic Design

- Usually a single large process
- Runs in a single address space – the kernel space
- Often a single binary file loaded at boot time
 - With other processes and files such as drivers assisting


Monolithic Design



Microkernel

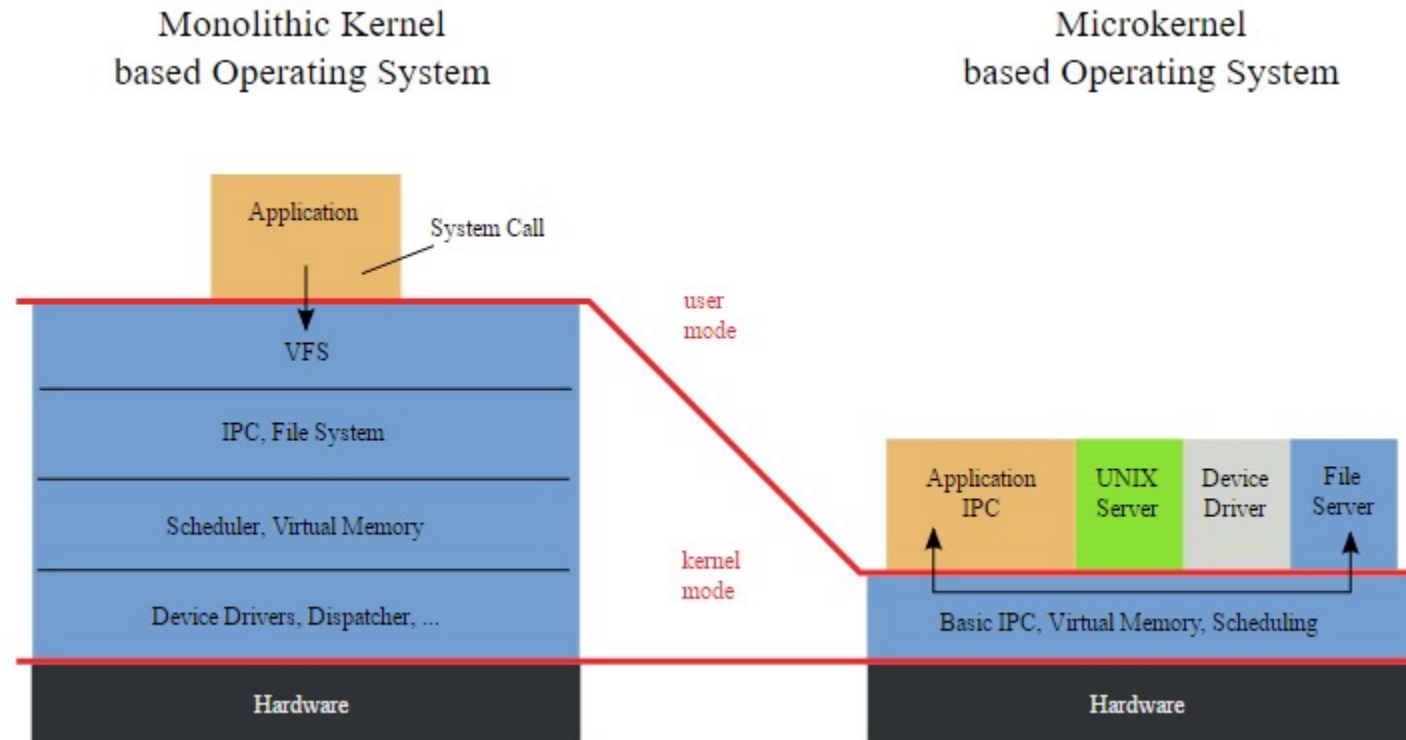
- Small kernel supporting minimal essential services
 - IPC, virtual memory, processing scheduling, interrupt handling, protection
- All other services handled by separate user level server processes
- Applications request for services from the server processes

Microkernel Design



Example: MINIX 3 – kernel only 12,000 lines

Microkernel – using IPC



Microkernel pros and cons

- Security advantage
 - Principle of Least Privilege
 - “only enough privilege to do what is required – and no more”
 - A problem bugging Windows
- Performance disadvantage
 - IPC requires 2 OS kernel crossing and process scheduling

Exokernels

- Squeeze kernel even further by removing more of its management duties
- Give applications direct access to hardware via libraries
- Experimental