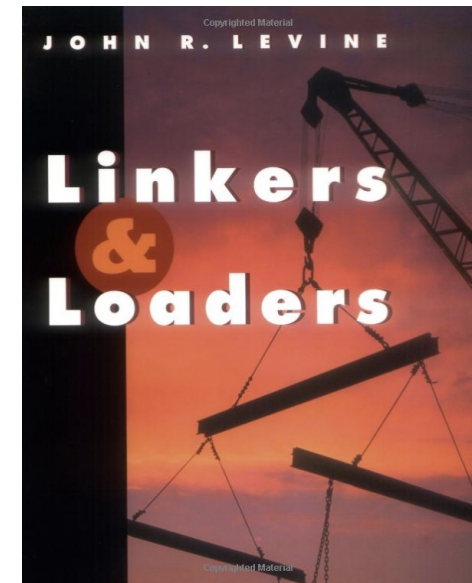


# Lecture 4

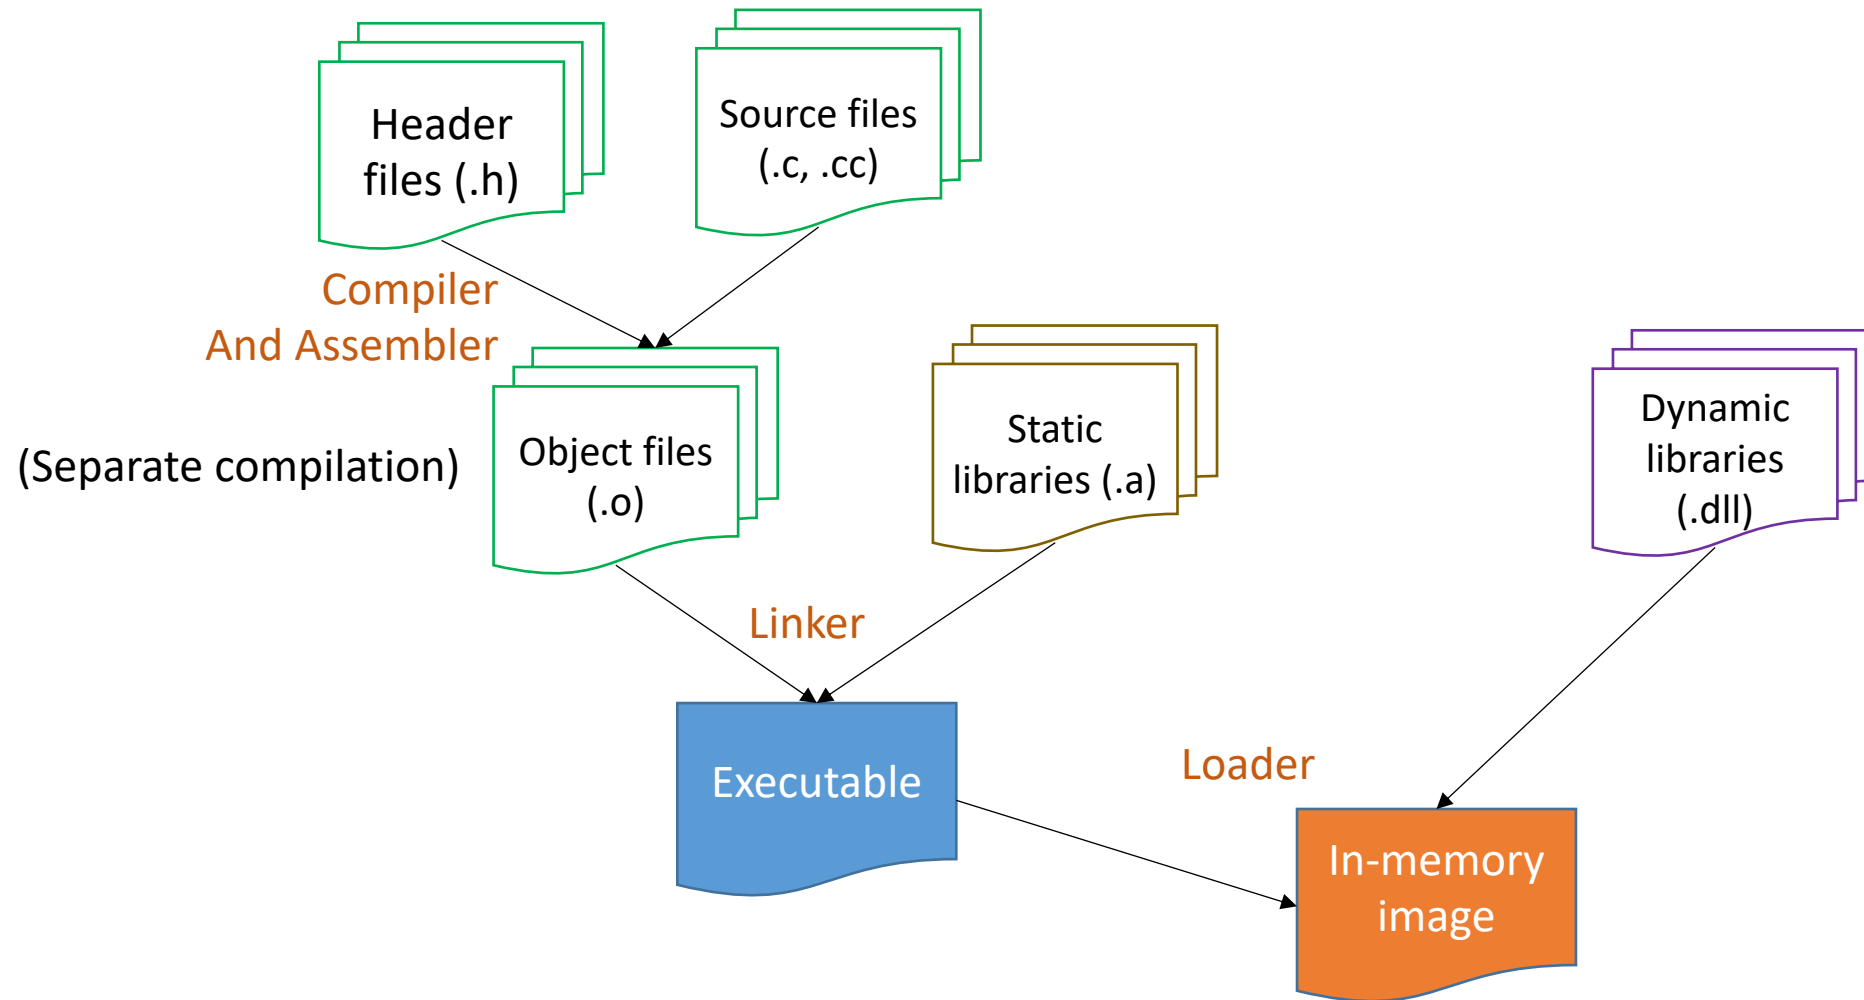
Linking and ELF  
(A small digression)

# Learning Objectives

- Learn about the linking and loader process
  - What are they and why they are designed in that way
- Learn the Executable and Linking Format (ELF)



# The compilation flow



# Main functions

- High level programming languages are good for human productivity
  - But they are not what the machine understands
- **Compilers**: translate code written in high-level programming languages to assembly code of the target machine's ISA
- **Assemblers**: translate human readable assembly code to binary object files
- **Linker**: combines multiple object files and library modules into a single executable
- **Loader**: reads the executable and together with dynamic libraries, constructs a memory image of the application

# Important design considerations

- Separate compilation
  - Allows for a better organization of the source code
    - Also facilitates sharing during code writing
  - Only need to recompile changed code
- External linkages
  - Modern applications are built on layers of libraries

# Executable files

- Information needs to be conveyed to the OS
- The OS expects executable files to have a specific format
  - *Header info*
    - Code locations
    - Data locations
  - Code & data
  - *Symbol Table*
    - List of names of things defined in your program and where they are defined
    - List of names of things defined elsewhere that are used by your program, and where they are used.

# Example – definition and use

```
#include <stdio.h>
```

```
int main () {
```

```
    printf ("hello, world\n")
```

```
}
```

- Symbol defined in your program and used elsewhere

- main

- Symbol defined elsewhere and used by your program

- printf

# A two-step operation

- **Linking**: Combining a set of programs, including library routines, to create a *loadable* image
  - a) Resolving symbols defined within the set
  - b) Listing symbols needing to be resolved by loader
- **Loading**: Starting from the loadable file, copy in dynamic libraries and construct the memory image in a new process



# Key activities involved

- **Relocation**: assigning load addresses to various parts of a program and adjusting the code and data in the program to reflect such assignments
- **Symbol resolution**: maps symbols into the actual location in a memory image
- **Loading**: constructing a memory image in a OS process from the executable on secondary storage

# Executable and Linking File Format

(ELF)

# ELF

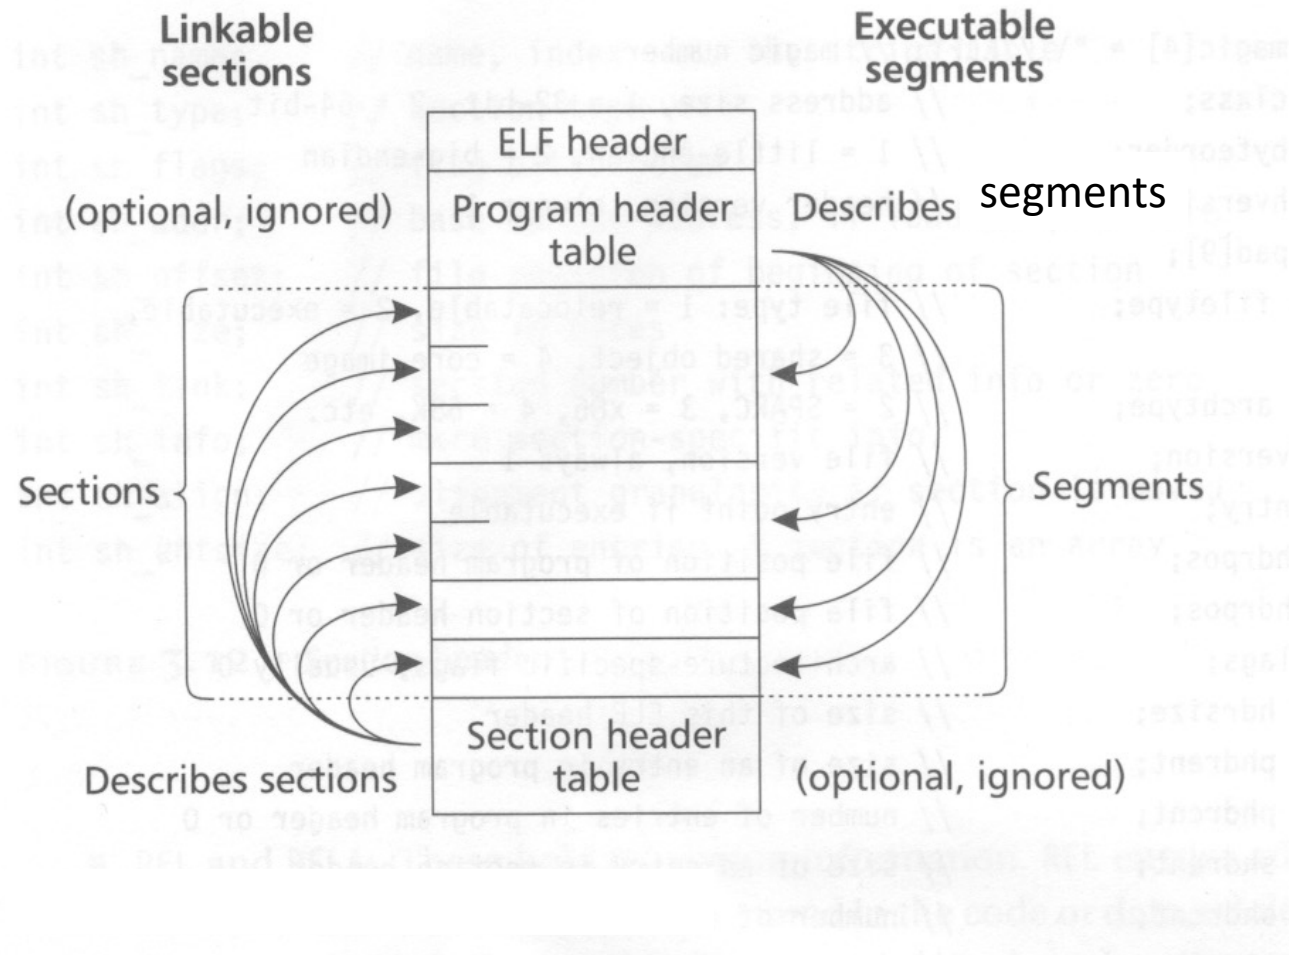
- Support for cross-compilation, dynamic linking, initializer/finalizer (e.g., the constructor and destructor in C++) and other advanced system features
- ELF has been adopted by FreeBSD and Linux as the current standard
- ELF32 for 32-bit binaries and ELF64 for 64-bit binaries
  - Almost identical except for length of data

# ELF File Types

- **Relocatable (object files)**
  - Created by compilers or assemblers. Need to be processed by the linker before running. Addresses assumed to start at zero.
- **Executable**
  - All relocation done and all symbol resolved except for shared library symbols that must be resolved at run time
- **Shared object**
  - Shared library containing both symbol information for the linker and directly runnable code
- **Core file**
  - For core dumps

# ELF Structure

- Compilers, assemblers, and linkers treat the file as a set of logical **sections** described by a section header table.
- The system loader treats the file as a set of **segments** described by a program header table.



# ELF Structure

- A single **segment** usually consist of **several sections**. E.g., a loadable read-only segment could contain sections for executable code, read-only data, and symbols for the dynamic linker.
- Relocatable files have section header tables. Executable files have program header tables. Shared object files have both.
- Sections are intended for further processing by a linker, while the segments are intended to be mapped into memory.
- See **`/usr/include/elf.h`**

# ELF Header

- The ELF header is always at offset zero of the file.
- The **program header table** and the **section header table**'s offset in the file are defined in the ELF header.
  - Use **fseek()** to find.
- The header is decodable even on machines with a different byte order from the file's target architecture.
  - After reading class and byteorder fields, the rest fields in the ELF header can be decoded.
  - ELF supports two different address sizes:
    - 32 bits
    - 64 bits



# ELF Program Header (32 bits)

```
typedef struct
{
    unsigned char e_ident[EI_NIDENT];    /* Magic number and other info */
    Elf32_Half    e_type;                 /* Object file type */
    Elf32_Half    e_machine;              /* Architecture */
    Elf32_Word    e_version;              /* Object file version */
    Elf32_Addr    e_entry;                /* Entry point virtual address */
    Elf32_Off     e_phoff;                /* Program header table file offset */
    Elf32_Off     e_shoff;                /* Section header table file offset */
    Elf32_Word    e_flags;                /* Processor-specific flags */
    Elf32_Half    e_ehsize;               /* ELF header size in bytes */
    Elf32_Half    e_phentsize;            /* Program header table entry size */
    Elf32_Half    e_phnum;                /* Program header table entry count */
    Elf32_Half    e_shentsize;            /* Section header table entry size */
    Elf32_Half    e_shnum;                /* Section header table entry count */
    Elf32_Half    e_shstrndx;             /* Section header string table index */
} Elf32_Ehdr;
```

4-byte magic number is “0x7F” followed by the string “ELF”.

# Relocatable Files

- A relocatable or shared object file is a collection of sections.
- Each section contains a single type of information, such as program code, read-only data, or read/write data, relocation entries, or symbols.
- Every symbol's address is defined relative to a section.
  - Therefore, a procedure's entry point is relative to the program code section that contains that procedure's code.

# Section Header

```
typedef struct
{
    Elf32_Word    sh_name;        /* Section name (string tbl index) */
    Elf32_Word    sh_type;        /* Section type */
    Elf32_Word    sh_flags;       /* Section flags */
    Elf32_Addr    sh_addr;        /* Section virtual addr at execution */
    Elf32_Off     sh_offset;      /* Section file offset */
    Elf32_Word    sh_size;        /* Section size in bytes */
    Elf32_Word    sh_link;        /* Link to another section */
    Elf32_Word    sh_info;        /* Additional section information */
    Elf32_Word    sh_addralign;   /* Section alignment */
    Elf32_Word    sh_entsize;     /* Entry size if section holds table */
} Elf32_Shdr;
```

# Types in Section Header (`sh_type`)

- **PROGBITS**: This holds program contents including code, data, and debugger information.
- **NOBITS**: Like PROGBITS. However, it occupies no space.
- **SYMTAB** and **DYNSYM**: These hold symbol table.
- **STRTAB**: This is a string table.
- **REL** and **RELA**: These hold relocation information.
- **DYNAMIC** and **HASH**: This holds information related to dynamic linking.

# Flags in Section Header (sh\_flags)

- **WRITE**: This section contains data that is writable during process execution.
- **ALLOC**: This section occupies memory during process execution.
- **EXECINSTR**: This section contains executable machine instructions.

# Sections

- `.text`:
  - This section holds executable instructions of a program.
  - Type: PROGBITS
  - Flags: ALLOC + EXECINSTR
- `.data`:
  - This section holds initialized data that contributes to the program's image.
  - Type: PROGBITS
  - Flags: ALLOC + WRITE

# Sections

- `.rodata`:
  - This section holds read-only data.
  - Type: PROGBITS
  - Flags: ALLOC
- `.bss` :
  - This section holds data with no initial values. The system will initialize the data to zero when the program begins to run.
  - Type: NOBITS
  - Flags: ALLOC + WRITE

# Sections

- `.rel.text`, `.rel.data`, and `.rel.rodata`:
  - These contain the relocation information for the corresponding text or data sections.
  - Type: REL
  - Flags: ALLOC is turned on if the file has a loadable segment that includes relocation.
- `.symtab`:
  - This section hold a symbol table.
- `.strtab`:
  - This section holds strings.



# Sections

- `.init:`
  - This section holds executable instructions that contribute to the process initialization code.
  - Type: PROGBITS
  - Flags: ALLOC + EXECINSTR
- `.fini:`
  - This section hold executable instructions that contribute to the process termination code.
  - Type: PROGBITS
  - Flags: ALLOC + EXECINSTR
- Programming language specific.
  - C does not need these two sections. However, C++ needs them.

# Sections

- `.interp`:
  - This section holds the pathname of a program interpreter.
  - Type: ALLOC
  - Flags: PROGBITS
  - If this section is present, rather than running the program directly, the system runs the interpreter and passes it the ELF file as an argument.
  - This facility runs non-text programs.
  - In practice, this is used to run the run-time dynamic linker to load the program and to link in any required shared libraries.
    - Variants of “`ld.so`” – the loader.

# Sections

- `.debug:`
  - This section holds symbolic debugging information.
  - Type: PROGBIT
- `.line:`
  - This section holds line number information for symbolic debugging, which describes the correspondence between the program source and the machine code (ever used gdb?)
  - Type: PROGBIT
- `.comment`
  - This section may store extra information.

# Sections

- `.got:`
  - This section holds the **global offset table**.
    - Crucial for shared library.
  - Type: PROGBIT
- `.plt:`
  - This section holds the **procedure linkage table**.
  - Type: PROGBIT
- `.note:`
  - This section contains some extra information.

# A typical relocatable file

ELF header	} (not considered sections)
(segment table)	
.text	
.data	
.rodata	
.bss	
.sym	
.rel.text	
.rel.data	
.rel.rodata	
.line	
.debug	
.strtab	
Section table	(not considered a section)

# String Table


- String table sections hold null-terminated character sequences, commonly called strings.
- The object file uses these strings to represent symbol and section names.
- We use an index into the string table section to reference a string.
- Separating symbol names from symbol tables frees us from any length limitation.

# Symbol Table

- An array that holds information needed to locate and relocate a program's symbolic definition and references.
- A **symbol table index** is a subscript into this array.

# Symbol Table

```
typedef struct
{
    Elf32_Word    st_name;           /* Symbol name (string tbl index) */
    Elf32_Addr    st_value;         /* Symbol value */
    Elf32_Word    st_size;          /* Symbol size */
    unsigned char st_info;          /* Symbol type and binding */
    unsigned char st_other;         /* Symbol visibility */
    Elf32_Section st_shndx;         /* Section index */
} Elf32_Sym;
```



The section relative in which the symbol is defined. (e.g., the function entry points are defined relative to .text)



# An example of a symbol table

```
[wongwf@localhost ~]$ cat t.c
#include <stdio.h>

int example_of_global_var = 0x123456;

main()
{
    printf("Hello world -- %d\n", example_of_global_var);
}
```

Symbol table '.symtab' contains 12 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	FILE	LOCAL	DEFAULT	ABS	t.c
2:	0000000000000000	0	SECTION	LOCAL	DEFAULT	1	
3:	0000000000000000	0	SECTION	LOCAL	DEFAULT	3	
4:	0000000000000000	0	SECTION	LOCAL	DEFAULT	4	
5:	0000000000000000	0	SECTION	LOCAL	DEFAULT	5	
6:	0000000000000000	0	SECTION	LOCAL	DEFAULT	7	
7:	0000000000000000	0	SECTION	LOCAL	DEFAULT	8	
8:	0000000000000000	0	SECTION	LOCAL	DEFAULT	6	
9:	0000000000000000	4	OBJECT	GLOBAL	DEFAULT	3	example_of_global_var
10:	0000000000000000	29	FUNC	GLOBAL	DEFAULT	1	main
11:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	printf

# .symtab vs .dynsym

- Sharable objects (libraries) and dynamic executables has two distinct symbol tables
  - $\text{.dynsym} \subseteq \text{.symtab}$
- **.symtab** is the full symbol table
  - Needed at static link time
  - Not all info needed at runtime hence not allocated in process memory
- **.dynsym** is the subset of .symtab needed at runtime
  - For runtime linking, loading and debugging

# Symbol Lookup

- Two ways to lookup:
  - If you know the index, then use index the symbol table
  - If you only have a name, then you need the `.hash` or `.gnu.hash` sections
    - If `symbol_table[hash_table[hash(symbol_name)]] == symbol_name` then a hit (ok, more complicated than this...)

# The four .gnu.hash sections

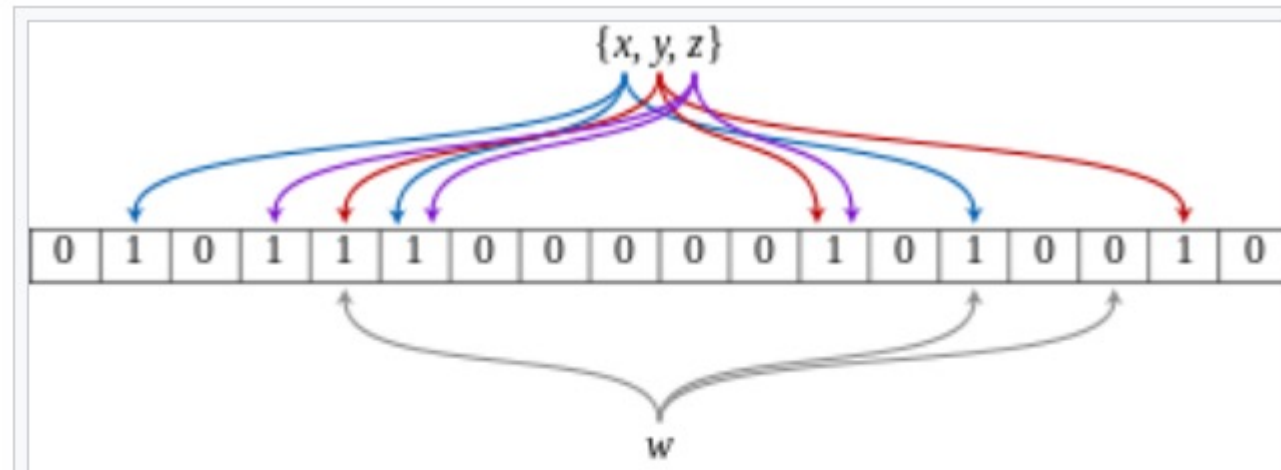
- Header: four 32-bit words
  - **nbuckets**: number of hash buckets
  - **symndx**: number of symbols of the dynamic symbol table that has been hashed
    - .dynsym may still contain other symbols not hashed
  - **maskwords**: number of words in the Bloom filter section
  - **shift2** – a shift count used in the Bloom filter
- Bloom Filter
- Hash Buckets
- Hash Values

# Bloom Filters

- A Bloom filter probabilistically tests whether an element is a member of a set.
- False positive matches are possible, but false negatives are not.
  - A query returns either “possibly in set” or “definitely not in set”.
- Elements can be added to the set, but not removed.
- The more elements that are added to the set, the larger the probability of false positives.

- Wikipedia

# Bloom Filter Example




An example of a Bloom filter, representing the set  $\{x, y, z\}$ .  
The colored arrows show the positions in the bit array that each set element is mapped to. The element  $w$  is not in the set  $\{x, y, z\}$ , because it hashes to one bit-array position containing 0. For this figure,  $m = 18$  and  $k = 3$ .

Wikipedia

# GNU Hash

- Uses k=2 Bloom filter

- $H1 = dl\_new\_hash(symbol\_name)$
- $H2 = H1 \gg shift2$

- $N = ((H1 / C) \% maskwords)$   
 C = size of one mask word in bits
- $BITMASK = (1 \ll (H1 \% C)) \mid (1 \ll (H2 \% C))$
- $bloom[N] \mid= BITMASK$
- **Test:**  $(bloom[N] \& BITMASK) == BITMASK$

```
static uint_fast32_t
dl_new_hash (const char *s)
{
    uint_fast32_t h = 5381;

    for (unsigned char c = *s; c != '\0'; c = *++s)
        h = h * 33 + c;

    return h & 0xffffffff;
}
```

# Hash buckets

- An array where each entry  $N$  is the lowest index into the dynamic symbol table for which:

$$(\text{dl\_new\_hash}(\text{symname}) \% \text{nbuckets}) == N$$

- **`dynsym[buckets[N]]`** is the first symbol in the hash chain that will contain the desired symbol if it exists.



# High level walkthrough - 1

- Compute **H1**, a 32 or 64 bit (depending on the ELF type) hash value for the symbol name (a C string)
- Use part of **H1** to be another hash **H2**
- Use H1 to compute **N**, an index into the Bloom array
- Create a BITMASK – where there is two 1's at  $(\mathbf{H1} \% \mathbf{C})$  and  $(\mathbf{H2} \% \mathbf{C})$
- Test if **bloom[N]** has two 1's at BITMASK positions
  - If false, we can confirm symbol not found

# Hash values

- Last of the 4 parts of .gnu.hash
- One entry for every (hashed) symbol of .dynsym.
- The top 31 bits of each entry contains the top 31 bits of the corresponding symbol's hash value.
- The least significant bit is used as a stopper bit.
  - It is set to 1 when a symbol is the last symbol in a given hash chain.

# Code walkthrough

Assume...

```
typedef struct {  
    const char    *os_dynstr;    /* Dynamic string table */  
    Sym           *os_dynsym;    /* Dynamic symbol table */  
    Word          os_nbuckets;   /* # hash buckets */  
    Word          os_symndx;     /* Index of 1st dynsym in hash */  
    Word          os_maskwords_bm; /* # Bloom filter words, minus 1 */  
    Word          os_shift2;     /* Bloom filter hash shift */  
    const BloomWord *os_bloom;   /* Bloom filter words */  
    const Word     *os_buckets;  /* Hash buckets */  
    const Word     *os_hashval;  /* Hash value array */  
} obj_state_t;
```

[https://blogs.oracle.com/ali/entry/gnu\\_hash\\_elf\\_sections](https://blogs.oracle.com/ali/entry/gnu_hash_elf_sections)

# Code walkthrough

```
Sym *
symhash(obj_state_t *os, const char *symname)
{
    Word      c;
    Word      h1, h2;
    Word      n;
    Word      bitmask;
    const Sym  *sym;
    Word      *hashval;

    /*
     * Hash the name, generate the "second" hash
     * from it for the Bloom filter.
     */
    h1 = dl_new_hash(symname);
    h2 = h1 >> os->os_shift2;

    /* Test against the Bloom filter */
    c = sizeof (BloomWord) * 8;
    n = (h1 / c) & os->os_maskwords_bm;
    bitmask = (1 << (h1 % c)) | (1 << (h2 % c));
    if ((os->os_bloom[n] & bitmask) != bitmask)
        return (NULL);
```

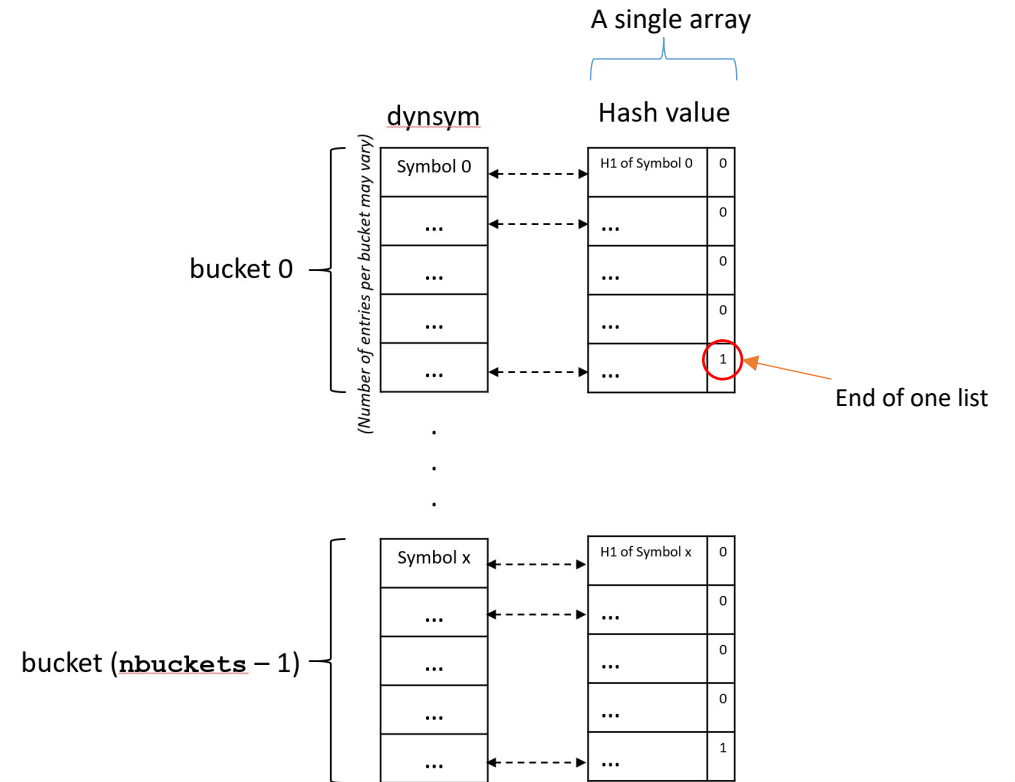
← If Bloom test fails, for sure not found.

## High level walkthrough - 2

- If it is true, we still need to check further coz Bloom filters can give false positives
- Dynamic symbol table organized by hash buckets
- Compute  $n$  – index into the hash bucket

# Dynamic symbol table hash buckets

- Scan list of hash value
- If match found, compare symbol name string



# Code walkthrough

```
/* Locate the hash chain, and corresponding hash value element */
n = os->os_buckets[h1 % os->os_nbuckets];
if (n == 0) /* Empty hash chain, symbol not present */
    return (NULL);
sym = &os->os_dynsym[n];
hashval = &os->os_hashval[n - os->os_symndx];

/*
 * Walk the chain until the symbol is found or
 * the chain is exhausted.
 */
for (h1 &= ~1; 1; sym++) {
    h2 = *hashval++;

    /*
     * Compare the strings to verify match. Note that
     * a given hash chain can contain different hash
     * values. We'd get the right result by comparing every
     * string, but comparing the hash values first lets us
     * screen obvious mismatches at very low cost and avoid
     * the relatively expensive string compare.
     */
    /* We are intentionally glossing over some things here:
     *
     * - We could test sym->st_name for 0, which indicates
     *   a NULL string, and avoid a strcmp() in that case.
     *
     * - The real runtime linker must also take symbol
     *   versioning into account. This is an orthogonal
     *   issue to hashing, and is left out of this
     *   example for simplicity.
     */
    /* A real implementation might test (h1 == (h2 & ~1), and then
     * call a (possibly inline) function to validate the rest.
     */
    if ((h1 == (h2 & ~1)) &&
        !strcmp(symname, os->os_dynstr + sym->st_name))
        return (sym);

    /* Done if at end of chain */
    if (h2 & 1)
        break;
}

/* This object does not have the desired symbol */
return (NULL);
```

# A small digression: Name Mangling

- Names in the symbol tables are often **not** the same names used in the source code
- Three reasons:
  - Avoid name collision
  - Name overloading
  - Type checking
- Unfortunately, no standard – every compiler does its own thing

```
$ c++filt _ZNK3MapI10StringName3RefI8GDScriptE10ComparatorIS0_E16DefaultAllocatorE3hasERKS0_  
Map<StringName, Ref<GDScript>, Comparator<StringName>, DefaultAllocator>::has(StringName const&) const
```

GNU C++ example - Wikipedia



# Relocation Table

- Relocation is the process of connecting symbolic references with symbolic definitions.
- Relocatable files must have information that describes how to modify their section contents.
- A relocation table consists on many relocation structures.
- Essentially “determine the value of X, and put that value into the binary at offset Y”

# Relocation Structure

- **r\_offset:**

- This field gives the location at which to apply the relocation.
- For a relocatable file, the value is the byte offset from the beginning of the section to the storage unit affected by the relocation.
- For an executable file and shared object, the value is the virtual address of the storage unit affected by the relocation.

# Relocation Structure

- **r\_info:**

- This field gives both the symbol table index with respect to which the relocation must be made and the type of relocation to apply.

- **r\_addend:**

- This field specifies a constant addend used to compute the value to be stored into the relocation field.

# Example

Relocation section '.rela.text' at offset 0x5d0 contains 3 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
0000000000000006	00090000000002	R_X86_64_PC32	0000000000000000	example_of_global_var - 4
000000000000000d	0005000000000a	R_X86_64_32	0000000000000000	.rodata + 0
0000000000000017	000b0000000002	R_X86_64_PC32	0000000000000000	printf - 4

Disassembly of section .text:

0000000000000000 <main>:

0:	55	push	%rbp	
1:	48 89 e5	mov	%rsp,%rbp	
4:	8b 05 00 00 00 00	mov	0x0(%rip),%eax	# a <main+0xa>
a:	89 c6	mov	%eax,%esi	
c:	bf 00 00 00 00	mov	\$0x0,%edi	
11:	b8 00 00 00 00	mov	\$0x0,%eax	
16:	e8 00 00 00 00	callq	1b <main+0x1b>	
1b:	5d	pop	%rbp	
1c:	c3	retq		

“In the final binary, patch the value at offset 0x6 in this object file with the address of symbol example\_of\_a\_global\_var”

# Executable Files

- An executable file usually has only a few segments. E.g.,
  - A read-only one for the code.
  - A read-only one for read-only data.
  - A read/write one for read/write data.
- All of the loadable sections are packed into the appropriate segments so that the system can map the file with just one or two operations.
  - E.g., If there is a `.init` and `.fini` sections, those sections will be put into the read-only text segment.

# Program Header

```
typedef struct
{
    Elf32_Word    p_type;           /* Segment type */
    Elf32_Off     p_offset;         /* Segment file offset */
    Elf32_Addr    p_vaddr;         /* Segment virtual address */
    Elf32_Addr    p_paddr;         /* Segment physical address */
    Elf32_Word    p_filesz;        /* Segment size in file */
    Elf32_Word    p_memsz;         /* Segment size in memory */
    Elf32_Word    p_flags;         /* Segment flags */
    Elf32_Word    p_align;         /* Segment alignment */
} Elf32_Phdr;
```

# The Types in Program Header

- **PT\_LOAD**: This segment is a loadable segment.
- **PT\_DYNAMIC**: This array element specifies dynamic linking information.
- **PT\_INTERP**: This element specified the location and size of a null-terminated path name to invoke as an interpreter.

# Executable File Example

	<i>File offset</i>	<i>Load address</i>	<i>Type header</i>
ELF header	0	0x80000000	
Program header	0x40	0x80000040	
Read-only text (size 0x4500)	0x100	0x8000100	LOAD, read/execute
Read/write data (file size 0x2200, memory size 0x3500)	0x4600	0x8005600	LOAD, read/write/ execute

Nonloadable information and optional section headers

FIGURE 3.16 • ELF loadable segments.



# An Example C Program

```
[wongwf@localhost ~]$ cat t.c
#include <stdio.h>

int example_of_global_var = 0x123456;

main()
{
    printf("Hello world -- %d\n", example_of_global_var);
}
```

# ELF Header Information

```
[wongwf@localhost ~]$ objdump -f a.out  
a.out:      file format elf64-x86-64  
architecture: i386:x86-64, flags 0x00000112:  
EXEC_P, HAS_SYMS, D_PAGED  
start address 0x0000000000400440
```

# Program Header

```
[wongwf@localhost ~]$ objdump -p a.out
```

```
a.out:      file format elf64-x86-64
```

```
Program Header:
```

PHDR	off	0x0000000000000040	vaddr	0x0000000000004000	paddr	0x0000000000004000	align	2**3
	filesz	0x00000000000001f8	memsz	0x00000000000001f8	flags	r-x		
INTERP	off	0x0000000000000238	vaddr	0x00000000000040238	paddr	0x00000000000040238	align	2**0
	filesz	0x000000000000001c	memsz	0x000000000000001c	flags	r--		
LOAD	off	0x0000000000000000	vaddr	0x00000000000040000	paddr	0x00000000000040000	align	2**21
	filesz	0x0000000000000071c	memsz	0x0000000000000071c	flags	r-x		
LOAD	off	0x00000000000000e10	vaddr	0x000000000000600e10	paddr	0x000000000000600e10	align	2**21
	filesz	0x00000000000000228	memsz	0x00000000000000230	flags	rw-		
DYNAMIC	off	0x00000000000000e28	vaddr	0x000000000000600e28	paddr	0x000000000000600e28	align	2**3
	filesz	0x000000000000001d0	memsz	0x000000000000001d0	flags	rw-		
NOTE	off	0x00000000000000254	vaddr	0x000000000000400254	paddr	0x000000000000400254	align	2**2
	filesz	0x00000000000000044	memsz	0x00000000000000044	flags	r--		
EH_FRAME	off	0x000000000000005f4	vaddr	0x0000000000004005f4	paddr	0x0000000000004005f4	align	2**2
	filesz	0x00000000000000034	memsz	0x00000000000000034	flags	r--		
STACK	off	0x00000000000000000	vaddr	0x00000000000000000	paddr	0x00000000000000000	align	2**4
	filesz	0x00000000000000000	memsz	0x00000000000000000	flags	rw-		
RELRO	off	0x00000000000000e10	vaddr	0x000000000000600e10	paddr	0x000000000000600e10	align	2**0
	filesz	0x000000000000001f0	memsz	0x000000000000001f0	flags	r--		

# Dynamic Section

```
Dynamic Section:
NEEDED               libc.so.6
INIT                 0x000000000004003e0
FINI                 0x000000000004005c4
INIT_ARRAY            0x00000000000600e10
INIT_ARRAYSZ          0x00000000000000008
FINI_ARRAY            0x00000000000600e18
FINI_ARRAYSZ          0x00000000000000008
GNU_HASH              0x00000000000400298
STRTAB                0x00000000000400318
SYMTAB                0x000000000004002b8
STRSZ                 0x0000000000000003f
SYMENT                0x00000000000000018
DEBUG                 0x00000000000000000
PLTGOT                0x00000000000601000
PLTRELSZ              0x00000000000000048
PLTREL                0x00000000000000007
JMPREL                0x00000000000400398
RELA                  0x00000000000400380
RELASZ                0x00000000000000018
RELAENT                0x00000000000000018
VERNEED               0x00000000000400360
VERNEEDNUM            0x00000000000000001
VERSYM                0x00000000000400358
```

```
Version References:
  required from libc.so.6:
    0x09691a75 0x00 02 GLIBC_2.2.5
```

Need to link this shared library  
for printf()

# Section Header

```
[wongwf@localhost ~]$ objdump -h a.out
a.out:      file format elf64-x86-64

Sections:
Idx Name          Size      VMA           LMA             File off  Algn
 0 .interp        0000001c  0000000000400238 0000000000400238 00000238 2**0
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 1 .note.ABI-tag   00000020  0000000000400254 0000000000400254 00000254 2**2
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 2 .note.gnu.build-id 00000024  0000000000400274 0000000000400274 00000274 2**2
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 3 .gnu.hash       0000001c  0000000000400298 0000000000400298 00000298 2**3
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 4 .dynsym         00000060  00000000004002b8 00000000004002b8 000002b8 2**3
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 5 .dynstr         0000003f  0000000000400318 0000000000400318 00000318 2**0
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 6 .gnu.version    00000008  0000000000400358 0000000000400358 00000358 2**1
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 7 .gnu.version_r  00000020  0000000000400360 0000000000400360 00000360 2**3
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 8 .rela.dyn       00000018  0000000000400380 0000000000400380 00000380 2**3
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 9 .rela.plt       00000048  0000000000400398 0000000000400398 00000398 2**3
   CONTENTS, ALLOC, LOAD, READONLY, DATA
10 .init           0000001a  00000000004003e0 00000000004003e0 000003e0 2**2
   CONTENTS, ALLOC, LOAD, READONLY, CODE
11 .plt            00000040  0000000000400400 0000000000400400 00000400 2**4
   CONTENTS, ALLOC, LOAD, READONLY, CODE
12 .text           00000184  0000000000400440 0000000000400440 00000440 2**4
   CONTENTS, ALLOC, LOAD, READONLY, CODE
13 .fini           00000009  00000000004005c4 00000000004005c4 000005c4 2**2
   CONTENTS, ALLOC, LOAD, READONLY, CODE
14 .rodata         00000023  00000000004005d0 00000000004005d0 000005d0 2**3
   CONTENTS, ALLOC, LOAD, READONLY, DATA
15 .eh_frame_hdr   00000034  00000000004005f4 00000000004005f4 000005f4 2**2
   CONTENTS, ALLOC, LOAD, READONLY, DATA
16 .eh_frame       000000f4  0000000000400628 0000000000400628 00000628 2**3
   CONTENTS, ALLOC, LOAD, READONLY, DATA
17 .init_array     00000008  0000000000600e10 0000000000600e10 00000e10 2**3
   CONTENTS, ALLOC, LOAD, DATA
18 .fini_array     00000008  0000000000600e18 0000000000600e18 00000e18 2**3
   CONTENTS, ALLOC, LOAD, DATA
19 .jcr            00000008  0000000000600e20 0000000000600e20 00000e20 2**3
   CONTENTS, ALLOC, LOAD, DATA
20 .dynamic        000001d0  0000000000600e28 0000000000600e28 00000e28 2**3
   CONTENTS, ALLOC, LOAD, DATA
21 .got            00000008  0000000000600ff8 0000000000600ff8 00000ff8 2**3
   CONTENTS, ALLOC, LOAD, DATA
22 .got.plt        00000030  0000000000601000 0000000000601000 00001000 2**3
   CONTENTS, ALLOC, LOAD, DATA
23 .data           00000008  0000000000601030 0000000000601030 00001030 2**2
   CONTENTS, ALLOC, LOAD, DATA
24 .bss            00000008  0000000000601038 0000000000601038 00001038 2**2
   ALLOC
25 .comment        0000002c  0000000000000000 0000000000000000 00001038 2**0
   CONTENTS, READONLY
```



# Symbol Table

```
[wongwf@localhost ~]$ objdump -t a.out
a.out:      file format elf64-x86-64

SYMBOL TABLE:
0000000000400238 l d .interp 0000000000000000 .interp
0000000000400254 l d .note.ABI-tag 0000000000000000 .note.ABI-tag
0000000000400274 l d .note.gnu.build-id 0000000000000000 .note.gnu.build-id
0000000000400298 l d .gnu.hash 0000000000000000 .gnu.hash
00000000004002b8 l d .dynsym 0000000000000000 .dynsym
0000000000400318 l d .dynstr 0000000000000000 .dynstr
0000000000400358 l d .gnu.version 0000000000000000 .gnu.version
0000000000400360 l d .gnu.version_r 0000000000000000 .gnu.version_r
0000000000400380 l d .rela.dyn 0000000000000000 .rela.dyn
0000000000400398 l d .rela.plt 0000000000000000 .rela.plt
00000000004003e0 l d .init 0000000000000000 .init
0000000000400400 l d .plt 0000000000000000 .plt
0000000000400440 l d .text 0000000000000000 .text
00000000004005c4 l d .fini 0000000000000000 .fini
00000000004005d0 l d .rodata 0000000000000000 .rodata
00000000004005f4 l d .eh_frame_hdr 0000000000000000 .eh_frame_hdr
0000000000400628 l d .eh_frame 0000000000000000 .eh_frame
0000000000600e10 l d .init_array 0000000000000000 .init_array
0000000000600e18 l d .fini_array 0000000000000000 .fini_array
0000000000600e20 l d .jcr 0000000000000000 .jcr
0000000000600e28 l d .dynamic 0000000000000000 .dynamic
0000000000600ff8 l d .got 0000000000000000 .got
0000000000601000 l d .got.plt 0000000000000000 .got.plt
0000000000601030 l d .data 0000000000000000 .data
0000000000601038 l d .bss 0000000000000000 .bss
0000000000000000 l d .comment 0000000000000000 .comment
0000000000000000 l df *ABS* 0000000000000000 crtstuff.c
0000000000600e20 l o .jcr 0000000000000000 __JCR_LIST__
0000000000400470 l F .text 0000000000000000 deregister_tm_clones
00000000004004a0 l F .text 0000000000000000 register_tm_clones
00000000004004e0 l F .text 0000000000000000 __do_global_ctors_aux
0000000000601038 l o .bss 0000000000000001 completed.6337
0000000000600e18 l o .fini_array 0000000000000000 __do_global_ctors_aux_fini_array_entry
0000000000400500 l F .text 0000000000000000 frame_dummy
0000000000600e10 l o .init_array 0000000000000000 __frame_dummy_init_array_entry
0000000000000000 l df *ABS* 0000000000000000 t.c
0000000000000000 l df *ABS* 0000000000000000 crtstuff.c
0000000000400718 l o .eh_frame 0000000000000000 __FRAME_END__
0000000000600e20 l o .jcr 0000000000000000 __JCR_END__
0000000000000000 l df *ABS* 0000000000000000
0000000000600e18 l o .init_array 0000000000000000 __init_array_end
0000000000600e28 l o .dynamic 0000000000000000 __DYNAMIC
0000000000600e10 l o .init_array 0000000000000000 __init_array_start
0000000000601000 l o .got.plt 0000000000000000 __GLOBAL_OFFSET_TABLE__
00000000004005c0 g F .text 0000000000000002 __libc_csu_fini
0000000000601034 g O .data 0000000000000004 example_of_global_var
0000000000000000 w *UND* 0000000000000000 __ITM_deregisterTMCloneTable
0000000000601030 w .data 0000000000000000 data_start
0000000000601038 g .data 0000000000000000 _edata
00000000004005c4 g F .fini 0000000000000000 __fini
0000000000000000 F *UND* 0000000000000000 printf@GLIBC_2.2.5
0000000000000000 F *UND* 0000000000000000 __libc_start_main@GLIBC_2.2.5
0000000000601030 g .data 0000000000000000 __data_start
0000000000000000 w *UND* 0000000000000000 __gmon_start__
00000000004005d8 g O .rodata 0000000000000000 .hidden __dso_handle
00000000004005d0 g O .rodata 0000000000000004 __io_stdin_used
0000000000400550 g F .text 0000000000000065 __libc_csu_init
0000000000601040 g .bss 0000000000000000 __end
0000000000400440 g F .text 0000000000000000 __start
0000000000601038 g .bss 0000000000000000 __bss_start
0000000000400530 g F .text 000000000000001d main
0000000000000000 w *UND* 0000000000000000 __Jv_RegisterClasses
0000000000601038 g O .data 0000000000000000 .hidden __TMC_END__
0000000000000000 w *UND* 0000000000000000 __ITM_registerTMCloneTable
00000000004003e0 g F .init 0000000000000000 __init
```

# Dynamic Symbol Table

```
[wongwf@localhost ~]$ objdump -T a.out
a.out:      file format elf64-x86-64

DYNAMIC SYMBOL TABLE:
0000000000000000      DF *UND* 0000000000000000 GLIBC_2.2.5 printf
0000000000000000      DF *UND* 0000000000000000 GLIBC_2.2.5 __libc_start_main
0000000000000000 w D *UND* 0000000000000000 __gmon_start__
```

# Dynamic Relocation Table

```
[wongwf@localhost ~]$ objdump -R a.out
a.out:      file format elf64-x86-64

DYNAMIC RELOCATION RECORDS
OFFSET          TYPE              VALUE
0000000000600ff8 R_X86_64_GLOB_DAT  __gmon_start__
0000000000601018 R_X86_64_JUMP_SLOT printf
0000000000601020 R_X86_64_JUMP_SLOT __libc_start_main
0000000000601028 R_X86_64_JUMP_SLOT __gmon_start__
```




# Allocation in Linking

# Need for allocation

- Each object file compiled differently with no idea about the others
- All objects in a given object file assumed to start at address 0
- Need to combine into a single executable

# Combining ELF object files

- Multiple segments needs to be combined
- Linker will collect the sections and place them into the segments of the final executable image plus write the program header



# You can control it!

- Linkers can be controlled via **link scripts**
- In GNU ld linker, you can use the “-T” option

```
SECTIONS
{
  . = 0x10000;
  .text : { *(.text) }
  . = 0x8000000;
  .data : { *(.data) }
  .bss : { *(.bss) }
}
```

“Start .text at virtual address 0x10000”

“Start .data followed by  
.bss at virtual address  
0x8000000”

A very simple example

# Dynamic Linking

# Motivation

- **Shared libraries:** only one copy of commonly used routines for all code in a system
  - Easy to maintain
  - Save space
- External symbols referenced in user code and defined in a shared library are resolved by the loader at load time.

# Advantages


- Faster load time: shared library code may already be in memory
- Better run-time performance: less likely to page out frequently used code
- Easier for shared libraries to be updated

# Disadvantages

- Need “glue code”
- Reduced locality of references
- May impact paging
- Changes in libraries may break some applications



# The Linux Library Hierarchy



# Static Libraries


- Simply a concatenation of object files with simple headers
- In Linux, it is created using the “**ar**” utility and has the file extension of “**.a**”

# Linux static library format

```
/* Archive files start with the ARMAG identifying string.  Then follows a
`struct ar_hdr', and as many bytes of member file data as its `ar_size'
member indicates, for each member file.  */

#define ARMAG    "<arch>\n"    /* String that begins an archive file.  */
#define SARMAG   8            /* Size of that string.  */
#define ARFMAG   "\n"         /* String in ar_fmag at end of each header.  */

struct ar_hdr
{
    char ar_name[16];          /* Member file name, sometimes / terminated.  */
    char ar_date[12];          /* File date, decimal seconds since Epoch.  */
    char ar_uid[6], ar_gid[6]; /* User and group IDs, in ASCII decimal.  */
    char ar_mode[8];           /* File mode, in ASCII octal.  */
    char ar_size[10];          /* File size, in ASCII decimal.  */
    char ar_fmag[2];           /* Always contains ARFMAG.  */
};
```



•  
•  
•

# Basically...

If `libx.a` consists of `obj1.o`, `obj2.o`, ..., `objk.o` concatenated

Then

```
gcc ... libx.a
```

Is just

```
gcc ... obj1.o obj2.o ... objk.o
```

# Dynamic linking: issues to resolve

- Allocate and load the segments found in the shared libraries
- Perform relocation
  - Recall: **relocations** are entries in binaries that are left to be filled in later
- Two solutions:
  - Load-time relocation
  - Position-independence code (PIC)

# Key Data Structures

- **Procedure Linkage Table (PLT)**
  - One entry for each external function reference.
  - All calls to the corresponding external function routed through this entry.
  - Easier to patch with actual address
- **Global Offset Table (GOT)**
  - All external global addresses.

# The special three entries of the GOT

- GOT[0]: address of the program's **.dynamic** segment
- GOT[1]: pointer to a linked list of nodes corresponding to the symbol tables for each shared library linked with the program
- GOT[2]: address of the symbol resolution function within the dynamic linker
  - **/lib{64}/ld-linux{-x86-64}.so**

# How it works (preparation - 1)

- Compiler ensures every call to the same dynamically linked routine will go to the same PLT entry
  - PLT[0] is reserved
- Each (64) bit PLT entry consists of:
  - **jmpq** \*GOT[entry corresponding to the function]
    - Called a “trampoline”.
  - **pushq** <GOT entry number of corresponding function>
    - Starts from 0, i.e. GOT[3].
  - **jmp** PLT[0]
- GOT[entry corresponding to the function] = address of “**pushq**”  
*initially*



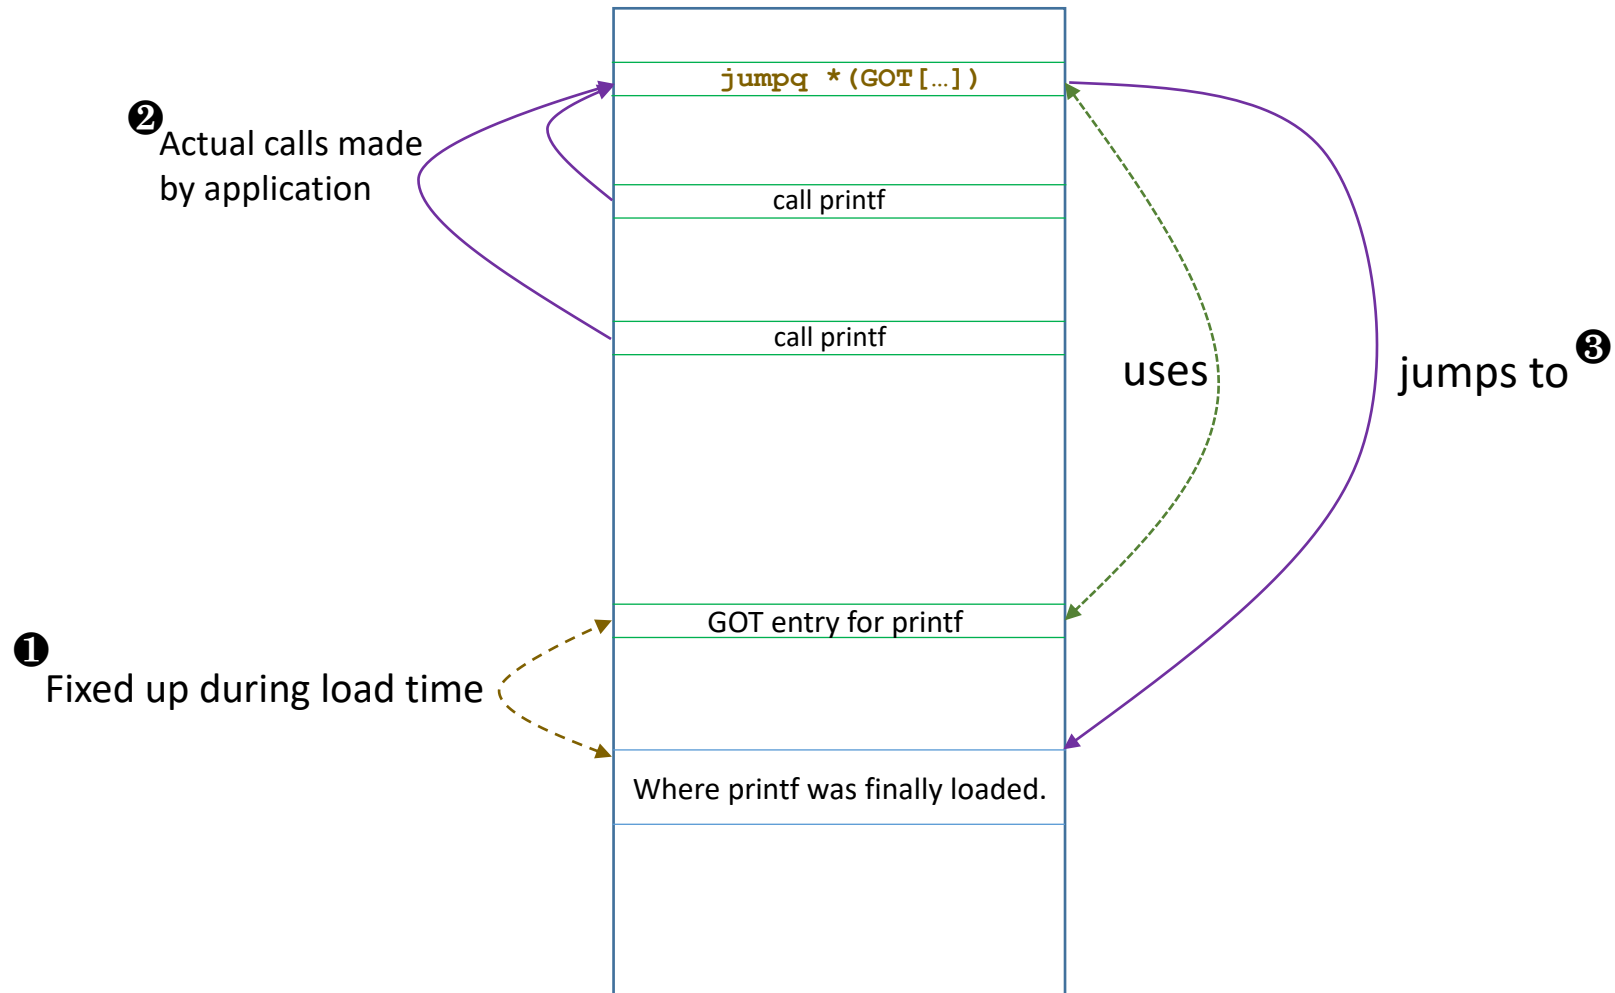
# How it works (preparation - 2)

- **.rela.plt** section of executable containing the GOT has the relocation information to fix up the GOT at runtime
- Address of **.rela.plt** itself is found as an entry in **.dynamic**
- Two ways to do dynamic linking:
  - “BIND NOW”
  - Lazy binding

# “BIND NOW”

- Find and load in all dynamic libraries
- Go through the main application's GOT relocation entries and fix all entries
- After fixing up, **GOT**[entry corresponding to the function] = address of actual function in the process image
- Therefore, trampoline jump **jmpq \*GOT**[entry corresponding to the function] will jump directly to the function
  - **pushq** and the second **jmp** never used

# “BIND NOW”



# Lazy Binding

- “BIND NOW” can slow down program start up
  - Define the environment variable **LD\_BIND\_NOW**
  - Many functions may not be used in a single execution anyway
- Solution: resolve binding only when we need it
- The reserved PLT[0] entry:
  - **pushq GOT[1]**
  - **jmpq \*(GOT[2])**

# GOT[2]

- Usually it is `_dl_runtime_resolve()`
  - In turn will call `_dl_fixup()`
- If profiling enabled, then `_dl_runtime_profile()`
- Assembly code in `glibc/dl-trampoline.S`

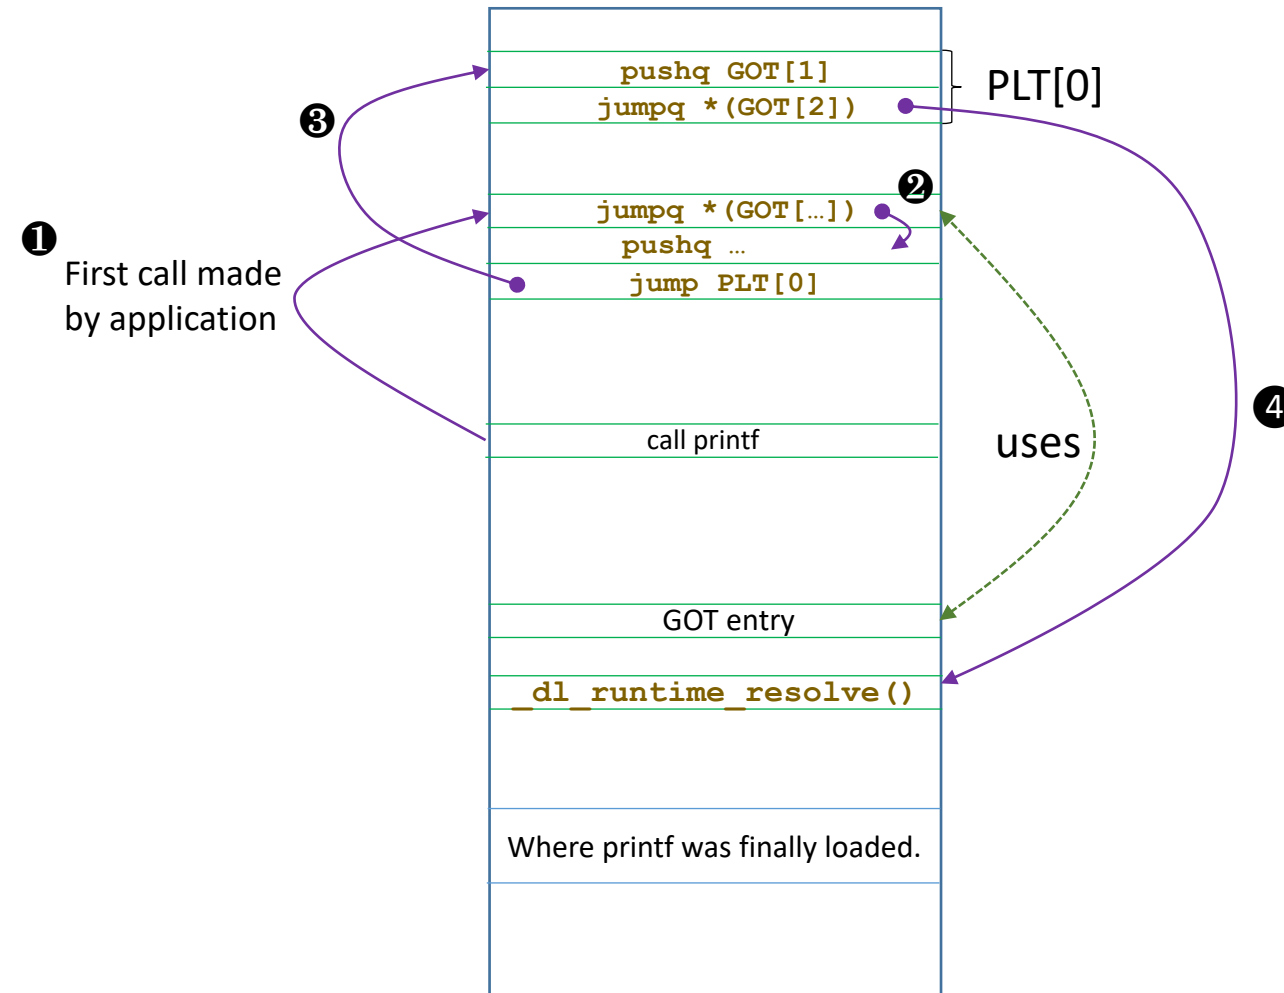
# Being lazy - 1

- Recall **GOT**[entry corresponding to the function] = address of “**pushq**” *initially*
- Lazy = not immediately change GOT entry at load time
- First time the PLT function entry is called, it will do **jumpq \* (...)** – jumps right back to the second instruction of the PLT entry, i.e. do:
  - **pushq** <GOT entry number of corresponding function>
  - **jmp** PLT[0]

## Being lazy - 2


- PLT[0] will jmp to `_dl_runtime_resolve()` of `/lib{64}/ld-linux{-x86-64}.so`
- `_dl_runtime_resolve()` will find final call site and using the relocation information patch the GOT entry
- Jumps back to PLT entry of the procedure. This time and all subsequent time, `jumpq * (...)` will jump using the patched GOT entry

# Being Lazy - 3





# Being Lazy - 4



# What about the shared objects themselves?

- Suppose a shared library itself uses a function in another?
- Shared libraries are also ELF
- They have their own PLT, GOT and relocation information embedded in the dynamic segment
- These too has to be resolved at runtime

# Position Independent Code

- Problem: hard to fix system-wide where a shared object should reside
  - Any one can create shared library.
  - Address Space Randomization (ASR)
- 64-bit Linux solution: shared objects must be compiled with “-fPIC”
- PIC will work anywhere it is placed
  - All branching uses relative addressing
  - Local data no issue (relative to stack or base pointer)
  - Global data: use indirect loading via GOT
  - Global shared procedure: use PLT and GOT

# (Just for fun) PIC example in 32 bit x86

```
#include <stdio.h>

int globvar = 42;

int foo(int arg)
{
    return globvar + arg;
}
```

C source

```
[CS5250]$ cc -S -m32 PIC-example.c
[CS5250]$ cat PIC-example.s
.file "PIC-example.c"
.globl globvar
.data
.align 4
.type globvar, @object
.size globvar, 4
globvar:
.long 42
.text
.globl foo
.type foo, @function
foo:
.LFB0:
.cfi_startproc
pushl %ebp
.cfi_def_cfa_offset 8
.cfi_offset 5, -8
movl %esp, %ebp
.cfi_def_cfa_register 5
movl globvar, %edx
movl 8(%ebp), %eax
addl %edx, %eax
popl %ebp
.cfi_restore 5
.cfi_def_cfa 4, 4
ret
.cfi_endproc
.LFE0:
.size foo, .-foo
.ident "GCC: (GNU) 4.8.5 20150623 (Red Hat 4.8.5-39)"
.section .note.GNU-stack,"",@progbits
```

```
[CS5250]$ cc -c -m32 PIC-example.c
[CS5250]$ objdump -d PIC-example.o
```

PIC-example.o: file format elf32-i386

Disassembly of section .text:

```
00000000 <foo>:
0: 55                push %ebp
1: 89 e5             mov %esp, %ebp
3: 8b 15 00 00 00 00 mov 0x0, %edx
9: 8b 45 08          mov 0x8(%ebp), %eax
c: 01 d0            add %edx, %eax
e: 5d               pop %ebp
f: c3               ret
```

Linker/Loader to fill this in

# (Just for fun) PIC example in 32 bit x86

```
[CS5250]$ cc -S -m32 -fPIC PIC-example.c
```

```
[CS5250]$ cat PIC-example.s
```

```
.file "PIC-example.c"
.globl globvar
.data
.align 4
.type globvar, @object
.size globvar, 4
```

```
globvar:
.long 42
.text
.globl foo
.type foo, @function
```

```
foo:
```

```
.LFB0:
```

```
.cfi_startproc
pushl %ebp
.cfi_def_cfa_offset 8
.cfi_offset 5, -8
movl %esp, %ebp
.cfi_def_cfa_register 5
call __x86.get_pc_thunk.cx
addl $_GLOBAL_OFFSET_TABLE_, %ecx
movl globvar@GOT(%ecx), %eax
movl (%eax), %edx
movl 8(%ebp), %eax
addl %edx, %eax
popl %ebp
.cfi_restore 5
.cfi_def_cfa 4, 4
ret
.cfi_endproc
```

```
.LFE0:
```

```
.size foo, .-foo
.section .text.__x86.get_pc_thunk.cx,"axG",@progbits,
.globl __x86.get_pc_thunk.cx
.hidden __x86.get_pc_thunk.cx
.type __x86.get_pc_thunk.cx, @function
```

```
__x86.get_pc_thunk.cx:
```

```
.LFB1:
```

```
.cfi_startproc
movl (%esp), %ecx
ret
.cfi_endproc
```

```
.LFE1:
```

```
.ident "GCC: (GNU) 4.8.5 20150623 (Red Hat 4.8.5-39)"
.section .note.GNU-stack,"",@progbits
```

%ecx contains EIP after call

Use GOT to indirectly obtain variable's address

Linker/loader to fill these in

Move return address to %ecx

```
[CS5250]$ objdump -d PIC-example.o
```

```
PIC-example.o: file format elf32-i386
```

Disassembly of section .text:

```
00000000 <foo>:
```

0:	55	push	%ebp
1:	89 e5	mov	%esp,%ebp
3:	e8 fc ff ff ff	call	4 <foo+0x4>
8:	81 c1 02 00 00 00	add	\$0x2,%ecx
e:	8b 81 00 00 00 00	mov	0x0(%ecx),%eax
14:	8b 10	mov	(%eax),%edx
16:	8b 45 08	mov	0x8(%ebp),%eax
19:	01 d0	add	%edx,%eax
1b:	5d	pop	%ebp
1c:	c3	ret	

Disassembly of section .text.\_\_x86.get\_pc\_thunk.cx:

```
00000000 <__x86.get_pc_thunk.cx>:
```

0:	8b 0c 24	mov	(%esp),%ecx
3:	c3	ret	

# (Just for fun) PIC made easier in x86-64

```
[CS5250]$ cc -S -fPIC PIC-example.c
[CS5250]$ cat PIC-example.s
.file "PIC-example.c"
.globl globvar
.data
.align 4
.type globvar, @object
.size globvar, 4
globvar:
.long 42
.text
.globl foo
.type foo, @function
foo:
.LFB0:
.cfi_startproc
pushq %rbp
.cfi_def_cfa_offset 16
.cfi_offset 6, -16
movq %rsp, %rbp
.cfi_def_cfa_register 6
movl %edi, -4(%rbp)
movq globvar@GOTPCREL(%rip), %rax
movl (%rax), %edx
movl -4(%rbp), %eax
addl %edx, %eax
popq %rbp
.cfi_def_cfa 7, 8
ret
.cfi_endproc
.LFE0:
.size foo, .-foo
.ident "GCC: (GNU) 4.8.5 20150623 (Red H
.section .note.GNU-stack,"",@progb
```



New  
addressing  
mode in  
x86-64

```
[CS5250]$ cc -c -fPIC PIC-example.c
[CS5250]$ objdump -d PIC-example.o

PIC-example.o:      file format elf64-x86-64
```

Disassembly of section .text:

```
0000000000000000 <foo>:
 0: 55                push    %rbp
 1: 48 89 e5          mov     %rsp,%rbp
 4: 89 7d fc          mov     %edi,-0x4(%rbp)
 7: 48 8b 05 00 00 00 00 mov     0x0(%rip),%rax
 e: 8b 10             mov     (%rax),%edx
10: 8b 45 fc          mov     -0x4(%rbp),%eax
13: 01 d0             add     %edx,%eax
15: 5d               pop     %rbp
16: c3               retq
```



Linker/Loader to fill this in

# Dynamic Loading

# Dynamic Loading API

- Loading a shared library under the control of the application
- Linux DL API
  - Must link with “**-ldl**” option

Function	Description
<b>dlopen</b>	Makes an object file accessible to a program
<b>dlsym</b>	Obtains the address of a symbol within a dlopened object file
<b>dlerror</b>	Returns a string error of the last error that occurred
<b>dlclose</b>	Closes an object file



# An example

```
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

int
main(int argc, char **argv)
{
    void *handle;
    double (*cosine)(double);
    char *error;

    handle = dlopen("libm.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(EXIT_FAILURE);
    }

    dlerror();    /* clear any existing error */

    /* Writing: cosine = (double (*)(double)) dlsym(handle, "cos");
       would seem more natural, but the C99 standard leaves
       casting from "void *" to a function pointer undefined.
       The assignment used below is the POSIX.1-2003 (Technical
       Corrigendum 1) workaround; see the Rationale for the
       POSIX specification of dlsym(). */

    *(void **) (&cosine) = dlsym(handle, "cos");

    if ((error = dlerror()) != NULL) {
        fprintf(stderr, "%s\n", error);
        exit(EXIT_FAILURE);
    }

    printf("%f\n", (*cosine)(2.0));
    dlclose(handle);
    exit(EXIT_SUCCESS);
}
```

Compile with:

```
gcc -rdynamic -o foo foo.c -ldl
```

End